

Flower Species Recognition on a Smartphone

Saahil Shihaz

Bachelor of Science in Computer Science
The University of Bath
2021-2022

This dissertation may be made available for consultation within the University Library and may be photocopied or lent to other libraries for the purposes of consultation.

Flower Species Recognition on a Smartphone

Submitted by: Saahil Shihaz

Copyright

Attention is drawn to the fact that copyright of this dissertation rests with its author. The Intellectual Property Rights of the products produced as part of the project belong to the author unless otherwise specified below, in accordance with the University of Bath's policy on intellectual property (see https://www.bath.ac.uk/publications/university-ordinances/attachments/Ordinances_1_October_2020.pdf).

This copy of the dissertation has been supplied on condition that anyone who consults it is understood to recognise that its copyright rests with its author and that no quotation from the dissertation and no information derived from it may be published without the prior written consent of the author.

Declaration

This dissertation is submitted to the University of Bath in accordance with the requirements of the degree of Bachelor of Science in the Department of Computer Science. No portion of the work in this dissertation has been submitted in support of an application for any other degree or qualification of this or any other university or institution of learning. Except where specifically acknowledged, it is the work of the author.

Abstract

The evolution of smartphones over time has made them more capable of complex tasks. One such task is image classification. This is where a computer attempts to identify an object or many objects within an image and give them a label. Classifiers can be built using machine learning (ML) by training models with example data so that it can attempt to recognise new data it hasn't seen before.

Deep learning (DL) is a part of ML that has seen massive growth in recent years. Recently, libraries like TensorFlow allow developers to implement DL models within smartphones. The state of the art will be highlighted within the project to provide a deep understanding of the context. Then an investigation is conducted to compare classification approaches for recognising flower species. The Inception V3 model is demonstrated to be very effective in identifying flowers within the Oxford Flowers 102 dataset, with an accuracy of 95.7%. An Android app will then be designed and developed using the latest available APIs to demonstrate the capabilities of the DL model within a smartphone by analysing how well the classifier performs.

Contents

1	Introduction	1
1.1	Problem Description	1
1.2	Main Objectives	3
2	Literature and Technology Review	4
2.1	Mobile Machine Learning	4
2.1.1	Evolution	5
2.1.2	Where does this lead us to, today?	6
2.2	Computer Vision	6
2.2.1	History	6
2.3	Machine Learning	7
2.3.1	Feature Extraction	8
2.3.2	Classification	8
2.4	Deep Learning	9
2.4.1	Neurons and Perceptrons	9
2.5	Convolutional Neural Networks	11
2.5.1	ML vs DL	12
2.6	Flower Classification	13
2.6.1	Existing Methods	13
2.6.2	Existing Apps	14
2.7	Evaluation	15
2.8	Summary	16
3	Investigation	17
3.1	Predictions	17
3.2	Design of Experiments	17
3.2.1	Oxford Flowers 102 Dataset	17
3.2.2	Classical Machine Learning	18
3.2.3	Deep Learning	19
3.2.4	Environment	20
3.2.5	Metrics	20
3.3	Results	21
3.3.1	Preliminary DL Findings	21
3.3.2	Performance	22
3.4	Analysis	23
3.4.1	Outcome	23
3.4.2	Comparison of Development Processes	24

3.4.3	Summary	25
4	Flower Classification App	26
4.1	Requirements and Specification	26
4.1.1	Literature and Technology Survey Findings	26
4.1.2	Google Developer Practices	27
4.1.3	Functional Requirements	28
4.1.4	Non-functional Requirements	28
4.2	Design and Implementation	29
4.2.1	Classifier	29
4.2.2	Implementation Process	29
4.2.3	User Interface Design	30
4.3	System Testing	31
4.3.1	Functionality Testing	31
4.3.2	Performance Testing	33
4.4	Summary	34
5	Conclusion	35
5.1	Contributions	35
5.2	Critical Appraisal	35
5.3	Future Work	36
5.4	Reflection	36
	Bibliography	38
A	Appendix	43
A.1	Table of Smartphones	43
A.2	Inception V3 Results Graphs	45
A.2.1	Full HP Results	45
A.2.2	Loss	45
A.3	Requirements	46
A.3.1	Functional Requirements	46
A.3.2	Non-functional Requirements	48
A.3.3	SDK percentage	48
A.3.4	Model performance list	49
A.4	Design	50
A.4.1	Prototype	50
A.5	Testing	50
A.5.1	Testing device specifications	50
A.5.2	Results of functionality testing	52
A.5.3	Results of distance testing	58
A.5.4	Results of lighting testing	61
A.5.5	Performance Timings	63
A.6	Google Image Classifier Example	65

List of Figures

2.1	Samsung Galaxy S21 Ultra 5G: Boasting an impressive array of camera sensors on the back (Three, 2021).	5
2.2	Simple diagram of a neuron (Scarpino, 2018).	10
2.3	Diagram of a perceptron (Anon, n.d.b).	10
2.4	Diagram of a layered network of perceptrons (Scarpino, 2018).	11
2.5	Example of a CNN process (Saha, 2018).	11
2.6	Diagram of PI@ntnet user scenario (Joly et al., 2016).	14
2.7	Screenshot of PictureThis app (PictureThisAI, n.d.).	15
3.1	Example images from the dataset.	18
3.2	The Inception V3 model diagram (GoogleCloud, 2022).	19
3.3	Graph generated in TensorBoard from HP training results.	21
3.4	Graph of accuracy against the number of training steps.	23
4.1	Flowchart that shows the actions that need to take place to classify a flower in the app.	29
4.2	Diagram that shows the processing the image has to go through until it gets fed to the classifier.	29
4.3	Sketch of the home page.	30
4.4	IDE preview of the home page.	31
A.1	Table of full HP training results as seen in TensorBoard.	45
A.2	Graph of loss against the number of training steps.	45
A.3	The percentage of devices that API level 8 will run according to Android Studio.	48
A.4	The list of models and their performance metrics on the TensorFlow site (TensorFlow, 2021).	49
A.5	Sketch of the flower information page	50
A.6	Classifying an Anthurium using the app.	53
A.7	Classifying a Columbine using the app.	54
A.8	Classifying an Orchid using the app.	54
A.9	Classifying a Daffodil using the app.	55
A.10	Classifying a Daisy using the app.	55
A.11	Classifying a yellow Dandelion using the app.	56
A.12	Classifying a white Dandelion using the app.	56
A.13	Classifying a Rose using the app.	57
A.14	Classifying a Rose at 4cm distance using the app.	58
A.15	Classifying a Rose at 15cm distance using the app.	59
A.16	Classifying a Rose at 35cm distance using the app.	59

A.17 Classifying a Rose at 60cm distance using the app.	60
A.18 Classifying a Rose in dark conditions using the app.	61
A.19 Classifying a Rose in normal conditions using the app.	62
A.20 Classifying a Rose in outdoor conditions using the app.	62
A.21 Execution timings of the CPU with 1 Thread.	63
A.22 Execution timings of the CPU with 2 Threads.	63
A.23 Execution timings of the CPU with 4 Threads.	64
A.24 Execution timings of the CPU with 8 Threads.	64
A.25 Execution timings of the GPU.	64
A.26 Output from the IDE when loading the example project.	65

Acknowledgements

I would like to thank my supervisor, Hongping Cai, for their valuable feedback and insight which helped improve the quality of the project. They provided me with great resources to help my understanding of the research topic and supported me through the entire year.

I would also like to thank my fellow coursemates for providing me with reassurance and support throughout the year. They really helped motivate me.

Finally, I would like to thank my friends and family as they would happily listen to me discuss my project and allow me to test my app on some of their flowers.

Chapter 1

Introduction

This section will outline the overall plan for this dissertation, starting with an in-depth look at the problem and a brief look at the domain.

1.1 Problem Description

The technological era that we live in has introduced many ground breaking achievements that constantly push the barrier of what is possible as well as introduce many new challenges that require complex solutions. One such challenge is big data processing, specifically, recognising patterns in data and drawing conclusions. Unfortunately machines don't have the ability to understand data the way that humans do and humans don't have the processing capability of modern machines. Due to obvious ethical and biological barriers, we cannot make humans fill the role of computers that compute data on a large scale, therefore, we must explore the alternative, making computers as smart as humans. This is where machine learning steps in, with which we have made great advancements in. What this project will focus on in particular, is granting the advanced capabilities of machine learning to lower end hardware.

This project aims to investigate the application of machine learning techniques to recognize images of flowers species on mobile devices. I will look at implementing standard machine learning algorithms that are effective in image classification as well as alternative deep learning techniques which are more effective in carrying out the same task. It will be interesting to compare both types of implementations in terms of accuracy, speed and performance and then transferring them to a mobile hardware environment which is traditionally weaker than standard machines such as desktops and laptops. Ultimately, we want to understand the best way in making a mobile software solution that can make use of machine learning methods and still maintain a seamless user experience.

Mobile devices have the advantage of portability and flexibility compared to PCs at the expense of pure processing power, storage and battery life. Advancements in machine learning can boost the abilities of mobile devices by allowing them to make informed decisions to aid the user. Traditional algorithms can't make decisions like "What is this flower?" without being cumbersome and inaccurate, we need something that can make good decisions and evolve, similar to human thinking. Smart phones and tablets are packed with more advanced technology than they've ever had like high resolution camera, sensors, displays and mobile processors, each of these are resources that a well written machine learning algorithm can take

advantage of, for example, in our case, a mobile phone that can provide high resolution photos of flowers. The more detailed data we can use to aid our machine learning process, the better.

Traditionally, mobile devices as well as similar devices with sensors do some light pre-processing of data, then they send it to the cloud which can handle actions that require intensive processing, this introduces some level of latency because of the communication between device and the cloud (Olascoaga, Meetr and Verhelst, 2021, p. 3). Latency, being a key issue, is important in some use cases such as autonomous vehicles, mobile gaming, activity tracking for vulnerable populations, etc (Olascoaga, Meetr and Verhelst, 2021, pp. 3-4).

With some raw information, a (classical) machine learning process can identify features, these could be used by a classifier that can make predictions given a set of data it hasn't seen before (LeCun, Bengio and Hinton, 2015). Features are sourced from the representation of an object, in turn the representation is defined by the data input. An example of a feature would be the presence or absence of thorns on the stem of a flower (Goodfellow, Bengio and Courville, 2016, p. 22). Traditional machine learning practices incorporated feature engineering that required designing custom algorithms for particular task which can be time consuming (Liu, Lin and Sun, 2020). There is also difficulty in understanding what features should be extracted, for example, it may be hard to represent flower petal shapes properly from raw pixel values if there are shadows being cast on it (Goodfellow, Bengio and Courville, 2016, p. 23). Representation learning is a method that can fix such issues by providing mappings not from just the representation of data to the output, but from representation to representation (Goodfellow, Bengio and Courville, 2016, p. 24). There are however, still hurdles to overcome, these are described as "factors of variation" where external factors might affect the source data, such as, the age of a flower, which could affect the petal shape and the season which may affect a flower's appearance. Factors like this make it difficult to get representations in the first place (Goodfellow, Bengio and Courville, 2016, p. 24).

Deep learning is a part of machine learning that aims to overcome limitations of classical machine learning techniques by expanding upon representation learning. Deep learning can be split into two unique parts:

- **Distributed Representation:** These are used to represent objects within a more compact and dense manner, instead of having representations for each type of object, for example a collection of words in a sentence, we could store the frequency of each word like the bag-of-words problem (Liu, Lin and Sun, 2020). This is a sparse representation and introduces problems with space and time complexity. Therefore distributed representations aim to tackle the sparsity problem as they are harder to model (Brownlee, 2017).
- **Deep Architecture:** The idea of layering to represent neurons in a human brain. You can imagine it as a map of nodes that takes an input, processes it through the different layers where at each step, a set of units calculate a weighted sum of their inputs from the previous layer and pass the result to the next layer until it gets to output units that generate a result. This is an example of a feedforward neural network (LeCun, Bengio and Hinton, 2015).

Input into a deep learning algorithm starts at the visible layer, which contains our set of input pixels that we can directly observe. This data is passed into a network of hidden layers, each of these layers represent an abstract feature that we can't normally observe by looking at the input data such as locations of edges and contours (Goodfellow, Bengio and Courville, 2016,

p. 26).

By making use of TensorFlow (Lite) we can produce classification models using languages like Python, C++ or Java, then convert said models into small packages that an Android/IOS application can use to generate predictions based on an input. TensorFlow is developed by Google and provides a suite of tools to design, test and deploy ML solutions. The Lite version that we will be using is designed specifically for mobile devices and IoT devices that may not have the support of powerful hardware. Using TensorFlow we can write models that use classical ML techniques as well as deep learning techniques like Convolutional Neural Networks (CNN) (Google, 2021b). There are examples that can be built specifically for flower classification within the API documentation which can serve as a starting point for the project.

1.2 Main Objectives

The main objectives serve as a plan to first, understand the domain, identify the key resources, experiment with approaches and then finally create a product that encapsulates everything that was learnt in this project. This is what I hope to achieve for this project:

- Identify and review literature that will aid in understanding the problem domain.
- Analyse existing mobile based image recognition software.
- Investigate the advantages and disadvantages of both classical ML and DL implementations and compare the two using analysis tools, this will be done on PC hardware.
- Design and implement an Android application that can recognize images using DL.
- Discuss the feasibility of DL on smartphones.
- Explore future improvements for the Android app.

Chapter 2

Literature and Technology Review

Ground-breaking achievements in technology and specifically machine learning have given us the tools and capabilities to tackle the key problem of image recognition. What this project aims to demonstrate is the application of deep learning within a mobile application to recognise flowers. It will be interesting to see how deep learning performs on mobile hardware which is generally less powerful than desktop PCs and laptops. Additionally, we will get to explore what sort of optimisations need to take place in order to get a feasible mobile based solution. The limited computing resources that we have to work with when creating our solutions is what makes this project challenging. Mobile devices typically contain smaller mobile processors, limited storage and batteries. These restrictions are in place to make mobile phones more efficient and to ensure that they last longer when not connected to a power source. In addition to this, the subject of machine and deep learning is complex and many find the concepts challenging to understand. What this literature review aims to do is identify key sources of information to help breaking down the underlying subject and discuss the quality of the research available.

2.1 Mobile Machine Learning

The potential of smartphones has still not been fully realised. With advancements in machine learning we can take the capabilities of smartphones to the next level by leveraging the advanced hardware within them to carry out complex tasks. This section will identify the key milestones within smartphone technology and how it leads us to integrating machine learning to make full use of the hardware.



Figure 2.1: Samsung Galaxy S21 Ultra 5G: Boasting an impressive array of camera sensors on the back (Three, 2021).

2.1.1 Evolution

Firstly, we want to look at the last 10 years of technological advancements within the smartphone space. With this information, we can hopefully gain some insight into the how much their capabilities have evolved. To do this, we will collect key aspects of specification data from the Samsung Galaxy flagship line of smartphones. Samsung currently hold the top spot in global market share at 20.8% as of Q3' 2021, this position is typically held by Apple or Samsung and can vary from a quarter to quarter basis (O'Dea, 2021).

By looking at the table in appendix A.1 we can see a clear increase in smartphone capability in multiple categories like the processor speed, core count, storage, memory and camera capabilities. There are of course many more different areas that are not listed like sensors, screen size and battery life which have also seen massive improvements over the last 10 years. We have only been seeing improvements in this space, therefore we can assume that we will continue to see improvements in the near future. What we must also consider is price and accessibility, flagship smartphones represent the top of the line offerings from each smartphone manufacturer and are of course priced as such. Low-end to mid-end smartphones are still more capable than their predecessors albeit their spec sheet may not be as impressive as high-end versions in the same generation. Therefore, we must ensure some level of scalability within our machine and deep learning processes. How can we make sure our processes can run efficiently on lower end hardware as well as high end hardware?

Kulendran et al. (2014), highlights how improvements in smartphones have created a boom in the number of smartphone based application designed to aid surgeons and patients in multiple facets of the medical industry like plastic, orthopaedic and general surgery. They conduct an expansive review of different solutions and analyse how the evolution of smartphones got them to the point that make them extremely useful as a tool to aid us. Ultimately, what this project aims to do is provide a robust software solution to recognise flower species on a smartphone, but we cannot ignore the fact that smartphones have come a long way in the hardware and operating system space to allow us to even conceive of a system.

2.1.2 Where does this lead us to, today?

ML and AI has become such an important part of smartphones that manufacturers now have dedicated processors for ML and AI tasks. Google includes a Tensor Processing Unit (TPU) in their Pixel line of phones (Triggs and Simons, 2021). Samsung, Qualcomm and Apple use their own solutions for machine learning processing by having their own bespoke processors. These processors are used to compute specific actions that require the decision making and accuracy capabilities of machine learning. Google Tensor in particular aids tasks such as speech recognition that is accurate but not taxing, therefore saving battery life. Tensor also applies to processing photographs and provides additional features to videos (Gupta, 2021). With such a focus on smartphones, to the point that they get dedicated hardware for ML, we should be seeing a huge increase in applications that integrate ML in some way, as well as the entire process of designing and implementing such solutions being carried out more rapidly, as developers learn to leverage the hardware.

2.2 Computer Vision

Since we are working with analysing images, the area of computer vision plays a big part in our research. In order to identify flower species we must first discuss techniques to analyse the incoming image data to make predictions using ML and DL.

2.2.1 History

Szeliski (2011) outlines significant occurrences in each decade starting from the 70s, thought to be the beginning of computer vision, all the way through to the 2000s. In the early 70s, researchers sought to emulate human intelligence in a machine by first solving the visual problem. It was hypothesised that if a computer could first recognize objects in the real world that it could then move onto the next step of using reasoning and problem solving at a high level. The first processes conducted to understand the 3D world were to extract edges to recognize 3D objects from 2D lines in an image.

The 80s were described to have a lot more focus on mathematical techniques for analysing scenes. Various algorithms and models were conceived as well as improvements in the contour and edge detection space. Researchers found that a lot of these algorithms could be thought of as “optimization problems” when they were described using the same mathematical framework.

We see more improvements in the field during the 90s including the production of 3D surfaces, tracking and image segmentation. However, what is probably more relevant to this project is statistical learning techniques that also started to appear during this decade. In 1991, we see a paper by Turk and Pentland (1991) that described the concept of “eigenfaces”. These are the product of converting images of faces into feature images. These feature images are essentially the training set. Recognition occurs “by projecting a new image into the sub-space spanned by the eigenfaces”. The new face is then classified by comparing its position relative to the known set of faces. Emphasis was placed on the limiting the scope of the allowed images, as such the system was trained and ready to accept profile straight-on images of the subject. In addition to that, they aimed to have the system compute a result in a reasonable time, which of course, is one of the goals of this project. The research hoped to improve on its predecessors that used, at the time, traditional methods of recognising features such as eyes’, noses and mouths and their relative position to each other. The work done with eigenfaces shows great similarities

with the machine learning techniques we see today, by essentially creating feature vectors and comparing the distance of known vectors in the same space.

Szeliski (2011) continues with their insight into the 2000s where we see the various improvements like more efficient algorithms and what finally dominates the latter half of the 2000s; applying machine learning techniques to computer vision to aid visual recognition research.

2.3 Machine Learning

This project will be using ML techniques to compare efficiency and accuracy to the more evolved deep learning. What we must first consider is how machine learning works in the context of computer vision. Camastra and Vinciarelli (n.d.) summarise this and break down ML development around three primary research points:

- Task-Oriented Studies, improving performance of learning systems in a predetermined set of tasks.
- Cognitive Simulation, emulating the human brain and designing processes around the human thought process.
- Theoretical Analysis, “the theoretical investigation of possible learning methods and algorithms independently of application domain”.

They also produce a taxonomy to represent the balance of two entities they describe: the “teacher” and the “learner”. The teacher, being the programmer, the one that designs the learning process and the learner being the computer system. The idea of inference is also introduced where a system can derive knowledge from previous observations. The taxonomy breaks down the amount of work that both the “learner” and the “teacher” need to do into four categories: Rote Learning, Learning from instruction, Learning by analogy and Learning from examples.

What we are more interested in is learning from examples where the “learner” infers the most out of the other categories in the taxonomy. The idea of the “learning problem” is introduced where the system needs to find a “general rule that explains the data given only a sample of limited size”. Learning techniques are broken down into four more categories: Supervised learning, Reinforcement learning, Unsupervised learning, Semi-supervised learning.

Zhu (2005) highlights semi-supervised learning in their survey as the combination of supervised and unsupervised learning where we use both labelled and unlabelled data for training of the classifier. They point to the survey done by Seeger (2000) in particular that provides more insight into the concept of semi-supervised learning. Their rationale for the concept in general was applying the ability for a system to make predictions based on knowledge it doesn't have. A supervised system has all labelled data to aid its training therefore its basis on making predictions is described as a “security belt” by Seeger. The model will basically make predictions within its limited scope, what we would call “overfitting” (Dietterich, 1995). Unsupervised learning heavily relies on prior assumptions for their final result this is because it doesn't have a knowledge base to rely on. By using a balanced combination of both implementations we can “balance the impact of prior assumptions”. Seeger also highlights the fact that labelling the data is a taxing process, fortunately for us, existing data sets already exist with labelled and unlabelled images for flower species which will be useful for training. Therefore, a semi-supervised approach is feasible for the ML approach of this project.

2.3.1 Feature Extraction

Agarwal (2021) provides an introductory guide to feature extraction. They describe feature extraction as one of the two ways to reduce dimensionality with the other being feature selection. Extraction produces new features which are described as “linear combination of the existing features”. The process aims to use less features to encapsulate the same image information.

Tian (2013) conducts a review of image feature extraction techniques that are worth considering. They start with discussing extracting colour features such as histograms and colour “moments” from specific colour spaces such as RGB and HSV. The paper also compares different types of colour features, for example, histograms are simple to compute but are sensitive to noise. This feature will be important as flowers come in many different colours, but it cannot solely be relied on as different species can share similar colours. We can also extract information about the texture of an image, this is where we start thinking about analysing groups of pixels together. Texture in the context of images is a way to describe the perceived smoothness, roughness or bumpiness of an image though spatial variations in pixel intensity levels (MathWorks, n.d.). Lastly, the paper goes into depth about shape features and points to different sources that go into the subject with more depth, but to summarise, shape features are split into two broad categories of contour and region based. This is where the features are calculated from shape boundaries and image regions respectively. A simple example of shape feature is the circularity ratio where you measure how close a shape is to a circle by calculating the ratio of the area of a shape to the area of a circle with the same perimeter (Mingqiang et al., 2008). Shape analysing will be very important in this project because flower shapes can differ greatly and can therefore serve as a way to easily differentiate between species.

2.3.2 Classification

Brownlee (2020a) provides an easy-to-understand breakdown of classification within ML. They describe it as the process of assigning “a class label to example from the problem domain”. In our case, that means classifying a flower as species A as opposed to B, C or D. They also go into detail about the different classification methods such as:

- Binary classification e.g. it's flower A or it's flower B.
- Multi-class classification, where we have more than two classes.
- Multi-label classification, this is where we have multiple predictions for classes based on a probability. This can be a path we could take if we wanted to produce multiple predictions for a flower species and then provide the likelihoods of each prediction to the user.

Next, we will look at classifiers, which help us carry out the classification stage. Fortunately, there is no shortage of types of classifiers in the ML space. Mohammed (2017) covers the most popular ones in good detail such as Naïve Bayes, k-Nearest Neighbour and Support Vector Machines (SVM). Starting with Naïve Bayes, this is a supervised classifier based on probability that assumes all attributes are independent:

$$P(c|E) = \frac{P(E|c)P(C)}{P(E)} \quad (2.1)$$

Where E is classified as the class $C = +$ if and only if

$$BC(E) = \frac{P(C = +|E)}{P(C = -|E)} \geq 1$$

BC is our Bayesian classifier, + and - are two separate classes (Zhang, 2004).

Zhang (2004) states that Naïve Bayes is superb in classification and demonstrates the classifier based version of it in Equation 1. They explore the optimal conditions of Naïve Bayes and propose that it is most optimal when the dependencies among attributes cancel out since Naïve Bayes works best when each attribute is independent, this is relevant as we want to explore optimisation in the project to ensure the highest level of efficiency when making the flower predictions.

Mohammed (2017) states that K-Nearest Neighbours (KNN) is one of the “simplest” of all the ML algorithms. Rosebrook (2016) discusses how to implement an image classifier using KNN where we can convert an image into a set of feature vectors on a graph, any new points get classified based on the k number of nearest neighbouring points, there's no real learning in this process, just the calculation of where the nearest points are, based on (usually Euclidean) distance.

Noble (2006) describes SVM as a way to tackle binary classifications, which means in the context of flower classification it answers questions like “is it Flower A or B?”. They state that you would need to train multiple “one-versus-all” classifiers to allow for multi-label classification.

ML techniques are certainly not useless and can still provide results, however, the research in the space has evolved to a new level, aiming to improve upon these traditional ML techniques in all evaluation categories. This is where Deep Learning (DL) comes in.

2.4 Deep Learning

DL will of course be our alternative approach to recognizing flower species. Where ML is basically our baseline, our DL implementation should hopefully highlight how much better it is compared to the ML approach.

2.4.1 Neurons and Perceptrons

Scarpino (2018) introduces the concept of Perceptrons in their book about using TensorFlow to implement DL. First, we must discuss neurons and how they relate to understanding the foundation of DL.

What they choose to highlight in particular are three points that describe a neurons functionality (and ultimately how it relates to perceptrons):

- A neuron receives one or more incoming signals and produces one outgoing signal.
- A neuron's output can serve as the input of another neuron.
- Every neuron has a threshold, and the neuron won't produce output until its electricity exceeds the threshold.

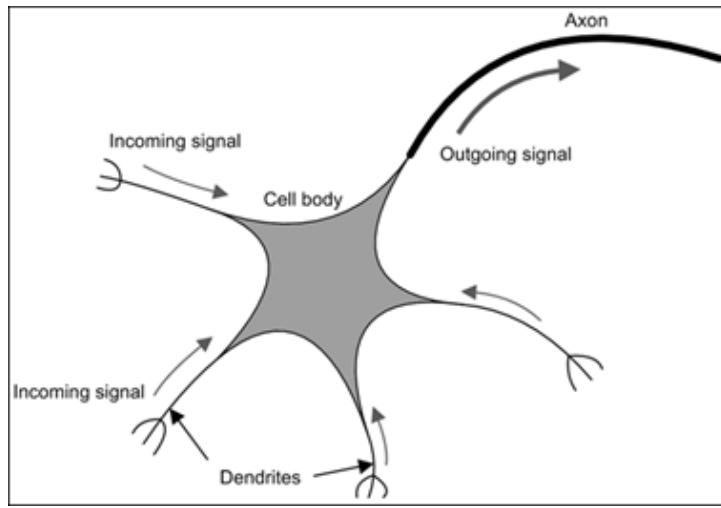


Figure 2.2: Simple diagram of a neuron (Scarpino, 2018).

This page by Anon (n.d.b) highlights a brief history of perceptrons, though it serves a starting point to learn more about the concept. Perceptrons were coined by Frank Rosenblatt in 1962 (Rosenblatt, 1961). His research is a bit outdated for our analysis, therefore here is a more modern version:

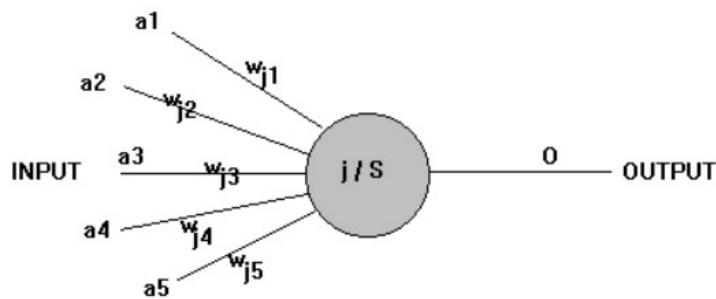


Figure 2.3: Diagram of a perceptron (Anon, n.d.b).

Each input on the left is weighted and summed within the circle node. If the summation meets a certain threshold, the output will be 1, if it doesn't meet the threshold then a 0 is outputted (Scarpino, 2018). Scarpino highlights some improvements to the model that was made including the weights that we discussed earlier as well as additional biases assigned with the incoming signals and an “activation function” that generates the output signal. Scarpino goes further by linking activation functions directly with in built TensorFlow functions that carry out the same task. Making it quite useful to understand the link between the TensorFlow API and the underlying DL context. Once we start linking perceptrons and arranging them into layers we get a neural network as shown here:

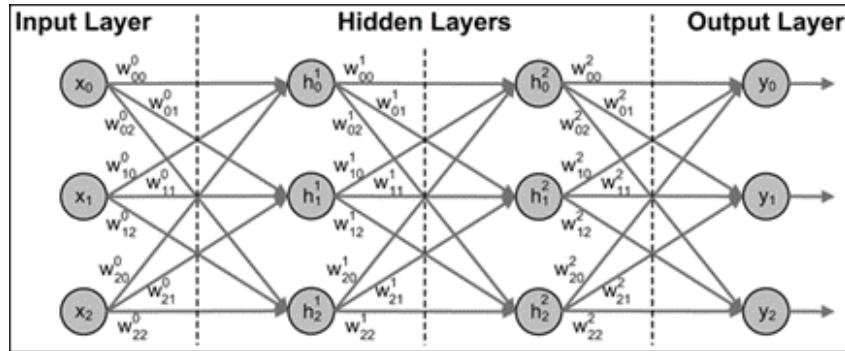


Figure 2.4: Diagram of a layered network of perceptrons (Scarpino, 2018).

2.5 Convolutional Neural Networks

Goodfellow, Bengio and Courville (2016) and Scarpino (2018) go into detail about CNNs, Scarpino in particular is a useful source on how it works with image classification in TensorFlow. However, it's useful to have some sort of starting point for the subject. Saha (2018) highlights the key features of a CNN and their purposes such as the individual layers: convolutional (kernel), pooling and classification (see Figure 2.5).

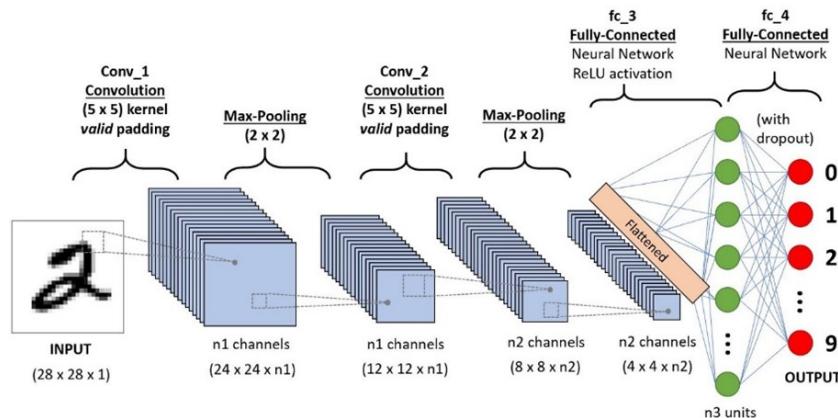


Figure 2.5: Example of a CNN process (Saha, 2018).

They also highlight a feature of CNNs that make images easier to process, where it reduces the size of images “into a form that is easier to process, without losing features which are critical for getting a good prediction”. This helps with our scalability approach when it comes to dataset sizes in particular. The ELI5 (Explain Like I'm 5) format is quite useful and allows us to highlight the key points of each layer (Saha, 2018):

- Convolution: Applying the kernel to extract high and low level features.
- Pooling: Reduces the spatial size of the output from the convolution. Decreases the “computational power” needed for data processing. Additionally, highlights features that are dominant.
- Classification: The pooling output is converted into column vectors and fed into a feed-forward neural network where the model is able to distinguish between features and classify them.

Saha (2018) also highlights that there are actually different implementations of CCNs, therefore they may not function in exactly the same format. Dive into Deep Learning (Anon, n.d.a) has a run down of multiple CNN types starting with LeNet-5 and more modern approaches like AlexNet, VGG, NiN, GoogLeNet, etc. We can also see how to implement them using TensorFlow which will prove useful when implementing our own solution for flowers.

Goodfellow, Bengio and Courville (2016) go into detail about Deep Learning from the concept of perceptrons to modern implementations. Something they talk about that is interesting is the increasing data set and model sizes over time, which is quite applicable to the project since we are working with modern mobile hardware. They discuss how the increasing capabilities of computer hardware have led to the development of larger models and that neural networks tend to double in size roughly every 2.4 years. They predict that the trend will continue further on in the future. What is also relevant is the dataset sizes. Storing datasets take up storage space, they state that a deep learning algorithm (as of 2015) is stated to perform at acceptable levels with “around 5,000 labelled examples per category, and will match or exceed human performance when trained with a dataset containing at least 10,000,000 labelled examples. That is of course, extremely large and is most definitely going to take up a lot of storage. Therefore, the book does re-iterate the earlier point of making use of unlabelled data like with semi-supervised learning. Goodfellow, Bengio and Courville (2016) dig deep into the subject of DL and explain subjects from the applied maths to the modern practices of DL and the research in the field. This can prove useful in fully understanding the various processes in place including optimisations and ways to increase accuracy that we will need to consider when designing a DL model for the mobile application.

2.5.1 ML vs DL

DL is an obvious evolution from ML, but it is worth highlighting the key differences for clarity because ultimately this project will compare how ML and DL compete with each other. Kavlakoglu (2020) breaks down how DL is different from ML. They highlight that DL takes the initiative by automating feature extraction to lessen human intervention and that ML is more reliant on humans, where humans normally define the characteristics to look out for, as well as their priorities. DL is stated to “require more data points to improve its accuracy” compared to the ML counterpart.

Xin et al. (2018) goes into depth about the key differences when discussing approaches to ML and DL in the context of cybersecurity. However, the same reasoning can be applied in our situation. The key points they highlight are:

- Data dependencies: DL performs better with larger datasets like mentioned earlier as well as ML outperforming DL with smaller data sets.
- Hardware dependencies: DL requires a lot of matrix calculations and therefore a Graphical Processing Unit (GPU) can be used to optimise these processes. Note that mobile hardware do contain GPU hardware but they are not on the same scale as dedicated GPUs you find in PC hardware. Therefore, it will be interesting to see how DL fares against ML when we keep this hardware dependency in mind.
- Feature processing: Once again iterating on the point mentioned before, DL can extract features directly from the data and requires less human intervention.
- Execution time: DL algorithms take a lot longer to train compared to ML, this is

dependant on the amount of data.

- Interpretability: Because of the complexity of DL it is hard to determine how a DL algorithm generated a result, whereas ML is more clearer.

I have summarised the key points, but they go into much more detail which could be helpful in the evaluation stage of comparing the two approaches of ML and DL.

2.6 Flower Classification

We will discuss further how flower classification is carried out including the use of ML and DL techniques, what features are extracted, the datasets that are used and the key challenges.

2.6.1 Existing Methods

Starting with what is known as the “Hello world” of ML, Iris flower classification serves as a simple and easy to understand project for developers to implement. The idea is to classify between three classes: Versicolor, Setosa and Virginica. There are many tutorials that can be followed online, this particular one by DataFlair (n.d.) provides additional background information about ML as well as how it will apply to the Iris project which is useful for our understanding. The tutorial uses the features of sepal length/width and petal length/width to determine the class of a flower. By using those inputs, they use a SVM to predict the species of a flower with 96% accuracy.

Nilsback and Zisserman (2008) demonstrate the effectiveness of a multiple kernel SVM on the oxford flowers 17 dataset. They manipulate the flower data to get key features such as the colour HSV values, the flower texture, shape and histogram of gradients (HOG) which “captures the more global spatial distribution of the flower” like the where the petals are arranged. They achieved an accuracy of around 88.3%. This is impressive considering the key challenges they highlight within flower classification. They state that flowers can share a lot of similarities between classes which can make it difficult to differentiate between species. Flowers are also “non-rigid objects” and therefore can appear in many different variations. Overall, they do a good job of explaining their reasoning for their dataset, citing the large variation of representations for each flower, and how they extract features from it, as well as how to build the classifier.

For deep learning, there is an extensive study that looks at using transfer learning, which is a technique of retraining TensorFlow models for different data sets. Xia, Xu and Nan (2017) use the Inception-v3 model to train a classifier for the Oxford-17 and Oxford-102 flower datasets. They go into detail about the steps that took place to carry out the transfer learning as well as how to reconfigure the last layer of the network to only have 17 and 102 outputs for each dataset as it defaults to 1000. They found that the model for the Oxford-17 and Oxford-102 datasets produced 95% and 94% accuracy respectively. This is very impressive performance, and their breakdown will be helpful when it comes to my own implementation. The paper really outlines the simplicity and flexibility of the Google’s TensorFlow, however, we will still need to investigate if these great results will translate to a mobile implementation as well.

2.6.2 Existing Apps

A large part of the project is to develop a fully functioning app that is developed to be more like a commercial product in addition to using deep learning techniques for the flower classification. This means developing features that will aid the user with using the app outside of main use case of recognising flowers. We will look at existing solutions that are already real products used by real people.

Pl@ntNet is a popular tool that has more than 10 million downloads on the Google Play Store alone (Google, 2022). It relies on volunteers to validate images and a search engine to identify them. Joly et al. (2016) go into detail about the overall experience of the app as well as provide insight into how it works.

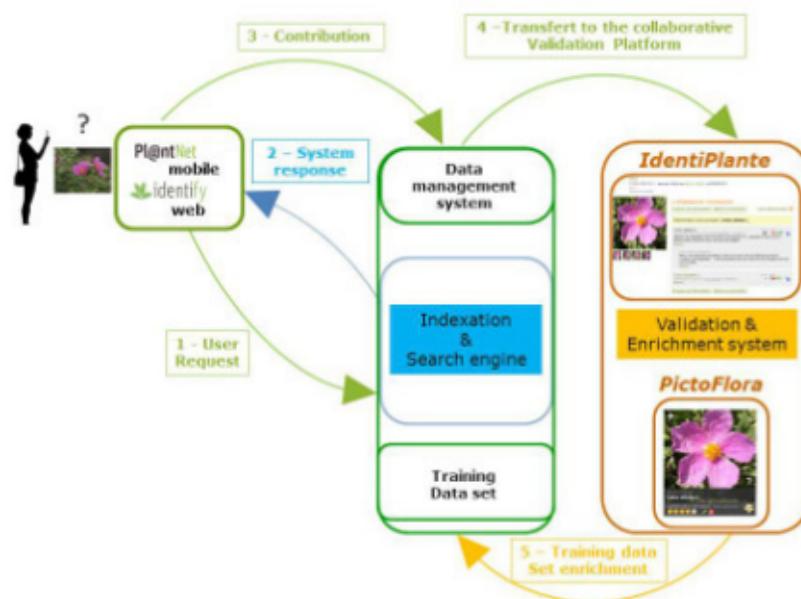


Figure 2.6: Diagram of Pl@ntnet user scenario (Joly et al., 2016).

Figure 2.6 shows clearly the type of system in place for the application. The user uses their device to query the search engine and get feedback, their image is also transferred to the collaboration platform if they chose to. It can then be independently verified and added back to the training set. The search engine is then retrained on a nightly basis. The paper doesn't go into much more detail about the search engine itself apart from mentioning that progress in machine learning and computer vision should improve the performance of identification. Unfortunately, there aren't any new papers that provide a better look at the app, so it is hard to understand what changes have been made over the last several years as well as how what methods they use to build the search engine. There is however, a dataset now available for use that covers over a thousand plant species with over 300 thousand images (Garcin et al., 2021). This is of course out of the scope of the project as it doesn't strictly contain flowers, but it does go into detail about how to use the dataset as well as how to load the data and build a model with it.

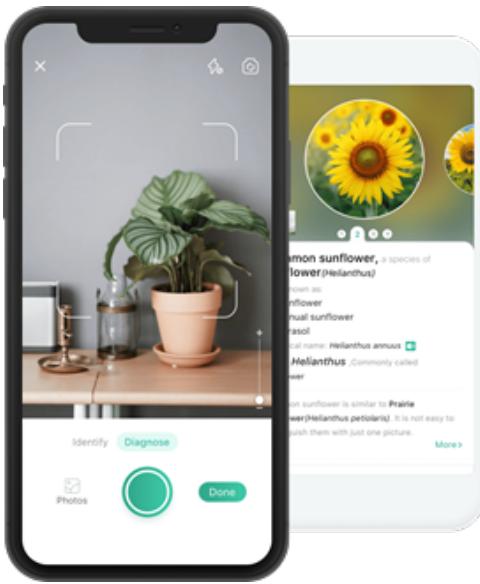


Figure 2.7: Screenshot of PictureThis app (PictureThisAI, n.d.).

PictureThis is another plant identification app (shown above) that doesn't go into detail about how it works but does note that it requires an internet connection to function properly. This suggests that it must communicate with a server in order to generate a prediction for images. The app is very streamlined and has an easy-to-use UI that also contains useful features such as how to care for the plant and important information about it. My approach will be different in the sense that any identification process will be carried out on the device, however, it is still important to consider alternative methods and how they perform, so that we can compare approaches.

2.7 Evaluation

One of the key points of the project is having to evaluate the ML and DL approach, what we haven't discussed yet however, is how do we go about doing this?

Williams, Zander and Armitage (2006) identify a process named "k-fold cross validation" where the data set is "divided into k subsets". One of the subsets is used to test the classifier and the rest ($k-1$) subsets is used as the training set. They use three performance metrics to test their ML systems: accuracy, precision and recall. Accuracy being the percentage of correct decisions over the total number of test instances. Confusion matrices can help us with representing accuracy by providing a "summary of prediction results" where we use the count of accurate and inaccurate predictions per class to show which particular classes the classifier may be struggling with (Brownlee, 2020b). Precision and recall are a bit more complex. Fortunately, Shung (2018) demonstrates how these two differ to accuracy. Precision is the number of instances that are correctly determined over the total number of instances that are guessed, this is made up of correctly guessed instances as well as instances that are incorrectly guessed. Recall is the number of correctly predicted instances over the true number of instances in the class. In addition to these evaluation methods, Williams, Zander and Armitage (2006) outline measuring CPU and memory usage. Fortunately, TensorFlow (lite) contains benchmarking tools for us to measure: Initialization time, inference time of warmup state/steady state, memory usage during initialization time and overall memory usage (Google, 2021a). I will also

assess real world speed and accuracy by analysing the app's performance during development. The TensorFlow (lite) guide also contains tutorials on how to choose the best model for the task by comparing model size and accuracy of different models as well as the time it takes to make the prediction. The lite version is designed specifically for mobile and internet of things (IoT) hardware, so the additional tools will prove useful for the project later when we are at the evaluation stage.

Chockwanich and Visoottiviseth (2019) use the same evaluation methods outlined when comparing different DL models implemented in TensorFlow. They also look at CPU usage percentages and processing time. They were able to make a clear conclusion of which model is better by evaluating all factors. A term called f1-score was also mentioned in their analysis. Shung (2018) also explains the relevancy of f1 score, it is essentially a way to determine a "balance between precision and recall". Korstanje (2021) discusses the F1 score and it's purpose in providing a better accuracy statistic that accounts for "imbalanced data", this is when you don't have a good balance of data for each class and therefore the classifier makes inaccurate predictions heavily skewed towards classes you have more data for. The F1 score is calculated by:

$$\frac{2 * Precision * Recall}{Precision + Recall}$$

2.8 Summary

The review has highlighted the progression ML and DL from the early concepts and how everything eventually fit together to form what we know today. The area will keep getting more exciting as we learn to optimise our current algorithms, come up with new ones and make use of advancing hardware capabilities. ML and DL is getting more and more accessible as manufacturers allow the use of their specialised hardware to developers who can make use of APIs built specifically for these tasks. Overall, we have built on the solid foundations of previous research and it's interesting to see how the field develops in the future. The project hopes to aid the field by investigating the ML and DL approaches in the mobile format as well as explore why we see certain results from the evaluation.

Chapter 3

Investigation

The main body of this dissertation will be split into two sections: the first being an analysis of TensorFlow's deep learning python API and how it compares to traditional machine learning python libraries like scikit. The second being the design and development of the flower classifier app. We will start with the investigation portion of the project, walking through my initial predictions, methodology and findings. The idea is to approach this investigation from the perspective of a software developer that is analysing the best approach for method of flower classification to use in their product. This includes assessing the quality of the resources available and discussing the possible challenges.

3.1 Predictions

The main prediction is that implementing a convolutional neural network is more suitable than using classical machine learning methods for this task. Suitability will be judged based on the metrics described in section 2.7 Evaluation. Furthermore, the process of implementing both approaches and the challenges that were faced will be described. Additional predictions mainly align with what was found in subsection 2.5.1 when discussing the key differences between ML and DL.

3.2 Design of Experiments

In this section, the implementation of the ML and DL classifiers will be individually described. The results from each approach will then be compared to fully understand the advantages of DL over ML. Furthermore, the development process including the various challenges faced will be discussed.

3.2.1 Oxford Flowers 102 Dataset

This dataset consists of 102 different flower species that occur within the UK. Each class contains between 40 and 258 images each. The images are described to have large variations between scale, pose and lighting, even within the classes itself. First, the dataset will be downloaded and managed by the TensorFlow Datasets module (TFDS) which will download the dataset to a generic directory and can directly manage the image and dataset data including filenames, class names and data splits. The data splits defined by TFDS consists of 6,149

images for the test set and 1,020 images for the train and validation sets each. This split is atypical as you would ideally have a larger number of images in the train set rather than the test set. The current split is thought to be a mistake on Google's side (TeaPearce, 2021). As a result, I will swap the train and test datasets and carry out the training process with the larger split which would be 75% of the dataset (TensorFlow, 2022a). Figure 3.1 below shows some of the example images in the dataset.

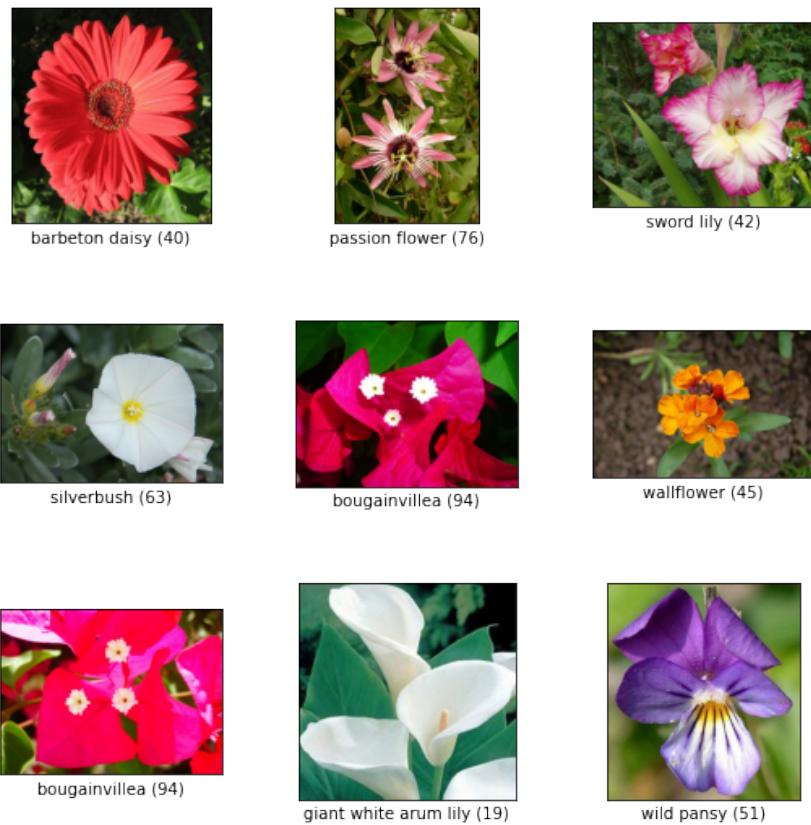


Figure 3.1: Example images from the dataset.

3.2.2 Classical Machine Learning

For this approach I decided to go with a Support Vector Machine (SVM) classifier that takes in features generated using “bag-of-words”, HSV (Hue, Saturation, Value) colour values and histogram of gradient values. I went with SVM because good performance has already been achieved with this particular dataset before with Nilsback and Zisserman (2008) as stated in the literature review. A value of 98.5% accuracy has also been achieved by using a CNN to extract the features (Mete and Ensari, 2019). I will not use a CNN to extract the features as the purpose of this investigation is to evaluate both approaches separately instead of combining aspects from both. Bag of words will be used as it is simple to implement, but will also provide consistent information about key points found in the image, despite how the image is presented in terms of factors like rotation and scale (Mohan, 2020). The “words” will be produced by extracting key points using Scale Invariant Feature Transform (SIFT) and then clustered by a K-Means trainer to make up the “vocabulary”. HSV values are useful as they give us the relevant colour data as well as information about the luminance in the image (Chapelle, Haffner and Vapnik, 1999). Histogram of gradients values provide information about the general shape

of the object; this is useful in mapping the various shapes and sizes that flowers come in. The dataset images will be resized to have a height and width of 299 pixels to match the input image conditions of the deep learning approach. All SVM parameters will be set to the default that is defined by the documentation.

3.2.3 Deep Learning

The CNN used in this approach will be Inception V3, a pre-trained network, specifically trained on the iNaturalist dataset, which contains 675,170 training and validation images from 5,089 categories (PapersWithCode, 2017). This means that the model has been optimised for recognizing plants and animals which makes it the best candidate to be used for transfer learning to allow it to recognize flower species. It is also listed as the second-best model hosted on TensorFlow. Inception V4 exists and has better performance to V3 but there are no fine-tuneable V4 models available that will allow transfer learning to take place.

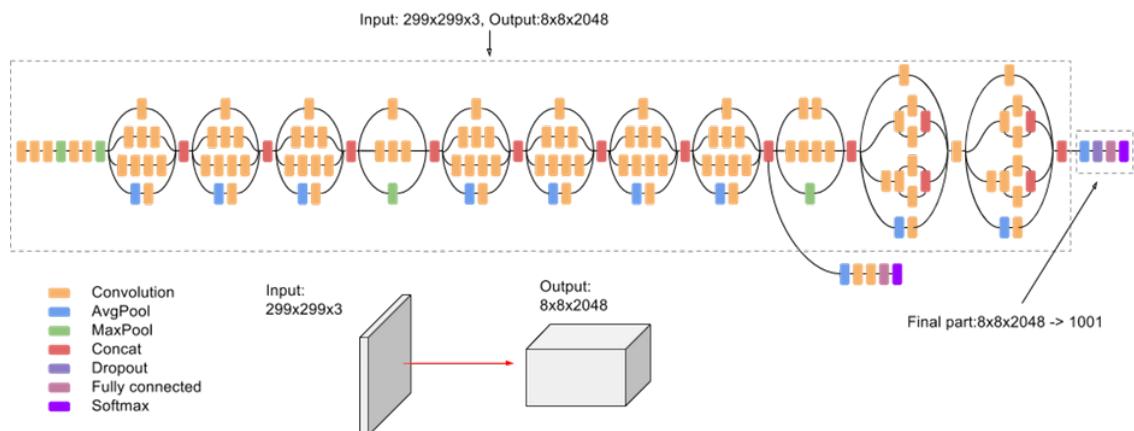


Figure 3.2: The Inception V3 model diagram (GoogleCloud, 2022).

Compared to the initial research of deep learning models presented in figure 2.5 within the literature review, figure 3.2 demonstrates how complex deep learning models can truly be. We discussed the terms convolution and pooling within section 2.5, however there are additional layers here that were not discussed:

- Concatenate (Concat) takes a list of tensors and outputs a single combined tensor (Keras, n.d.b).
- Fully connected layers map all inputs in the previous layer to every “activation unit” of the layer next to it (Pratap Singh, n.d.).
- Softmax is where probabilities are assigned to each label based on the likelihood of the image belonging to that label. All probabilities add up to 1 (GoogleDevelopers, 2020d).
- Dropout mitigates overfitting of a dataset by randomly dropping out neurons on each pass of the network when training (Sebastian, 2021).

The images will go through some data pre-processing such as resizing to have a width and height of 299 as that is the requirement of the input tensor for Inception V3. They will also have to have their values rescaled from 0-255 to between 0-1. Random flips and crops will also be added as part of the pre-processing pipeline. Images will be batched into sets of 32

images, this means that 32 images will be trained per step per epoch. Hyper parameter tuning is also required to try and get the best performance possible. These are the hyper parameters will be adjusted and their supposed effect:

- **Optimiser** is used to improve the speed and performance of the model by adjusting the parameters of the model during training to minimise loss and maximise accuracy (Maithani, 2021). The types of optimiser that will be tested are Adam, Stochastic Gradient Descent (SGD) and AdaMax. There are a few more optimisers available that are not listed, as they are unsuitable for this dataset.
- **Learning rate** is the rate at which a model learns, a larger value means that the model learns faster at the expense of producing substandard weights for the model (Perlato, n.d.). Values from 0.01-0.0001 will be tested, moving down a magnitude at each step.
- **Dropout** which was described earlier when discussing the Inception V3 model. Increasing this value will mean a larger percentage of nodes will get removed. Values within the range 0.2-0.4, in increments of 0.1 will be tested.

Using a TensorFlow module called TensorBoard, we can test the hyper parameter combinations efficiently and produce the metrics for each combination. TensorBoard allows the developer to view how the hyper parameters affected the results. In an effort to decrease overall training time and save time, I will conduct some preliminary testing to see which optimiser is more suitable with just baseline parameters. Once that is selected, I only need to test the different learning and dropout rates using a grid search. This is when every possible combination is tested (REFERENCE HERE). When testing with a large number of hyper parameters, one can use other methods like random search to decrease overall tuning time by randomly sampling hyper parameters from a range based on a statistical distribution, this means that more effective hyper parameters are tested to avoid spending time on hyper parameters that will not affect the overall performance that much (Sayak, 2018).

3.2.4 Environment

Both approaches will be developed and ran on the same machine, a custom desktop PC that contains these main components:

- An AMD Ryzen 3600 4.2Ghz 6 Core/12 Thread CPU
- 16GB 3200Mhz DDR4 Memory

The PC ran on Ubuntu 20.04 LTS with the relevant python3 and TensorFlow libraries needed to run the Jupyter Notebooks locally. VS Code was used as the Integrated Development Environment (IDE) with the Python and Jupyter extensions. It is possible to run model training on the GPU instead of the CPU, however, only Nvidia GPUs are directly supported. As a result of not having access to one, we will be mainly using the power of the CPU to carry out training.

3.2.5 Metrics

The key metrics I will analyse along with accuracy, precision, recall and F1 are:

- Loss, a measure of how bad the model's predictions are (GoogleDevelopers, 2020c). This metric only applies to the deep learning model as it is calculated during the training

process as the model tries to minimise it.

- Area under the receiving operating characteristic curve (AOC). You get an ROC curve by plotting the true positive rate against the false positive rate. The area under this curve indicates how well the model is selecting the correct prediction against all other predictions (GoogleDevelopers, 2020b).

3.3 Results

3.3.1 Preliminary DL Findings

Another factor to consider when carrying out training is the number of epochs. An epoch is a full pass over the training set (REFERENCE HERE). Multiple passes are needed to minimise loss and fully train the model. Through preliminary testing of the model, I found that 5 epochs are more than sufficient as we reach the maximum validation accuracy shown in figure 3.4. If we increase the number of epochs, we risk overfitting (GeeksForGeeks, 2020). Optimisers were tested individually to determine how much they would affect the results under standard conditions. I found that SGD was unsuitable for this task, producing accuracy results at almost half of what Adam and AmaMax were producing. Overall, Adam produced the best results, with AdaMax being only a percentage point behind. Therefore, I decided to conduct any main hyper parameter tuning, purely using the Adam optimiser. All initial training is done to only the first epoch to reduce overall execution time.




Figure 3.3: Graph generated in TensorBoard from HP training results.

Figure 3.3 above is a parallel co-ordinates view shown within TensorBoard that clearly show what accuracy results certain combinations of hyperparameters produce. Highlighted in green shows the joint first highest with a dropout of 0.3 and a learning rate of 0.001. A dropout of 0.2 with the same learning rate produces the same result. A learning rate of 1e-4 produces a significant reduction of first epoch performance, this is because we would need to increase the number of overall epochs to get optimal results. Ideally, we would test the number of epochs along with the other hyperparameters to produce fairer results, however, that would significantly increase training time with a grid search. If a developer has access to more specialised hardware, suited for machine learning tasks, this would not be too much of an issue.

Overall, the final model will use a learning rate of 0.001 and dropout of 0.3. This was decided after considering the results from the hyper parameter tuning (Appendix A.2.1) as well as the proposed learning rate being the default one provided by the TensorFlow API.

3.3.2 Performance

Here are the overall performance metrics produced from both approaches:

Metric	SVM (%)	Inception V3 (%)
Accuracy	24.30	95.69
Loss	-	17.95
Precision	29.40	96.14
Recall	24.30	95.69
F1	26.60	95.91
ROC AUC	89.60	99.97

Table 3.1: Results output from the classifier after predicting against the testing dataset.

Precision, recall and ROC AUC values are weighted, meaning that they consider class balance as they calculate the metrics for each class and find the average weighted by the number of correct instances for each class (Scikit, n.d.). This was necessary to account for the slight imbalances we have with the number of images in each class. The deep learning approach clearly comes out on top with excellent results.

The Inception V3 test contains additional information about accuracy (figure 3.4) and loss (Appendix A.2.2) for the training and validation datasets while the model is being trained. It tells us how the model improves per epoch.




Figure 3.4: Graph of accuracy against the number of training steps.

The model needs to be converted to a TensorFlow lite file in order for it to be allowed to be used in the Android application. Once we convert it, we can reload the model and run inference on it to check if it is still effective. The result after doing that produces an accuracy of 100% if we just test it on the first batch of 32 images. Of course, this is not indicative of real-world performance, and we will need to do additional real-world performance profiling once the app is developed.

3.4 Analysis

3.4.1 Outcome

It is clear that the DL approach is vastly superior in classifying flowers than the SVM in all aspects. The Inception V3 approach has high accuracy meaning it can correctly identify the classes of most of the images in the test set. High precision indicates that a large portion of correct identification were genuinely correct. High recall demonstrates how well the classifier identified correct instances. You might notice that accuracy and recall are the same value for both classifiers, this is because they represent the same thing in non-binary classification. Recall is shown to be:

$$\frac{TP}{TP + FN} \quad (3.1)$$

Where TP is number of true positives and FP is the number of false negatives (GoogleDevelopers, 2020a).

In this case, FN is every incorrect evaluation of a class. Therefore, if you calculate this recall value for each class, you're essentially tallying up the number of times the class has been correctly identified over the total number of images in the class for every class, giving us the total accuracy. $F1$ gives us a better indication of the accuracy of the classifier, by taking into account the precision and recall, the high $F1$ score indicates that even with class imbalance, we get good classification results. An AUC of near 100% means that the predictions are nearly always correct. The other predictions of an image are taken into account, so while the classifier might choose the incorrect overall label, if the true label has a high probability, this means that the classifier is still somewhat good at identifying the correct label.

The SVM falls behind in every aspect, therefore, it is not worth pursuing as a method for flower classification in this dataset. It is still possible to adjust the parameters of the SVM to get better performance or get more feature data from the feature extraction methods. The current implementation, specifically the histogram calculations, use parameters that decrease extraction and training time significantly, in exchange for outputting less information. Nilsback and Zisserman (2008) achieved an accuracy of 88.3% with the oxford flowers 17 dataset, a smaller subset of the one being used here using very similar methods of combining SIFT, HOG and HSV and using an SVM. Their method is a lot more intuitive than my SVM implementation but still doesn't reach the performance of the Inception V3 model. This doesn't mean that SVM or any other classical machine learning method is obsolete, even in being used as a classifier on mobile phones. They can still perform extremely well on smaller and less complex datasets. Therefore, if dataset complexity is not an issue, what else is there to consider? We can also look at the ease of implementation and how easy it is to move the classifier into a mobile phone.

3.4.2 Comparison of Development Processes

Most of the SVM was implemented with the help of scikit which contains easy to understand documentation on how to implement an SVM as well as many other machine learning algorithms. There is also example code that gives developers an idea on how to implement the algorithms in Python. Scikit provides a more extensive support library for hyper parameter tuning compared to TensorFlow, including in-built functions for grid search. TensorFlow uses TensorBoard as a platform to display the results of tuning as well as the overall results of training, which is very useful, but doesn't provide any additional API level support for finding the right hyper parameters. That is up to the developer to implement themselves. Scikit provides additional tools to allow model persistence so that developer trained models can be saved and reused later. However, allowing these saved models to be used within an android application doesn't appear to be straight forward compared to the TensorFlow approach and there does not seem to be any documentation to help with this.

The documentation for TensorFlow is extremely robust. This is unsurprising, considering it is developed and maintained by Google. There is a clear-cut development process for the developer, where a beginner, can understand the basics and go onto develop and train their first neural network on any of the supported platforms. The Keras Python API runs on top of the TensorFlow platform and is designed to be easy for developers to use and produce results quickly (Keras, n.d.a). My approach to this involved adapting the example code provided in the TensorFlow API. This allowed me to quickly get a working example that I could adjust by

simply searching through the API documentation. The documentation is easy to follow as it provides clear explanations as well as examples on how it would be used. The process to train models for different datasets using transfer learning is simple as well, allowing developers to adapt any model. We can then save the model to storage so that it can be reloaded, and re-training is not required. The saved model can then be converted to be used as a TensorFlow Lite model to be used in a mobile phone app. The process for this is simple and requires using the example code provided to implement. The developer needs to adjust and provide the details to describe the model, including the shape of the input tensor, the image data type, and the number of classes. This data can then be used by the TensorFlow Lite interpreter within Android.

Due to time and hardware restrictions, tuning the hyper parameters of the deep learning models had to be short and straight-forward. Given better hardware resources such as dedicated Nvidia GPUs and larger memory capacities, it would've been possible to conduct more efficient hyper parameter tuning that could lead to better model performance.

3.4.3 Summary

This investigation has been very educational and has certainly shown the capabilities of these specially made software libraries. It is worth going back to discuss the main differences of ML and DL discussed in section 2.5.1 and see if they still hold. Data dependencies are first discussed, and it is certainly impressive how well Inception V3 handles the Oxford flowers 102 dataset compared to the SVM. The model is capable of handling datasets that contain up to 1001 classes, therefore, it is possible to push the model even further. It would've been interesting to see how effective a GPU would be with training so that comparisons could be made on how much of an improvement it is over just using a CPU. The DL approach certainly required less set up and thought when it came to actually implementing a classifier. This is because the model can directly extract features and the developer doesn't have to worry about what features to extract like they do with an SVM. The training time is about the same for both approaches. When increasing the number of features produced from the bag of words, HSV and HOG extractors, the training time would also increase significantly. Similarly, with the DL model, when decreasing the learning rate and adding more epochs, the training time will increase. Therefore, that comparison is situational. When it comes to understanding the functionality of the two approaches, it was certainly a requirement to understand the SVM approach relatively in depth compared to the DL approach. This is mainly because of the feature extraction process, where developers can clearly see what exactly is being analysed and the exact format of the input data. Whereas, with DL, it appears to be more like a black box, where the complexity of the model is abstracted away from the developer. Therefore, the only requirements were to provide the data in the format expected for the model to function properly and you would get the results you need. TensorFlow has shown itself to be a very capable library that allows developers to easily implement deep learning for many use cases and I am excited to see how it improves in the future.

Chapter 4

Flower Classification App

The second section of this project involves designing, implementing and testing an Android app that can recognise flower species. I will proceed with using the TensorFlow library to seamlessly integrate the deep learning model developed in section one into an app. I will discuss the various processes that will take place, challenges faced and real-world performance of the final product.

4.1 Requirements and Specification

Requirements were defined based on the findings produced in the literature review. This includes the background research related to the goals of deep learning and the findings from the existing solutions review. I will also use the guidelines produced by Google when it comes to TensorFlow integration, Material You design and Android development. These guidelines are useful in designing and developing applications that conform to the general Google standards. As we don't have a defined client or target market, there are no requirements based on questionnaires and interviews. Therefore, my aim is to describe the development process and test the performance of the app so that the feasibility of using deep learning within a mobile phone can be assessed.

4.1.1 Literature and Technology Survey Findings

The Literature and Technology review walked through the evolution of the domain and how DL was the product of many years of research and development. It is important, to re-iterate the reason why there is a large number of resources dedicated in the progress of this space. The goal is to imitate human-like decision making within a computer. In the context of flowers, this means the ability for a computer to recognise different flower species accurately and efficiently. This means that the system must be able to recognise the flowers correctly and in a timely manner.

PI@ntNet is the most robust existing solution I have reviewed with a system that can identify a large number of different plant species. Due to the many different categories of plant life that the app supports, the user interface uses additional settings to narrow the search. The process of taking a picture and then getting the results is quite simple, however they are carried out within separate views. This means that taking a picture, viewing the picture that was taken afterwards and then getting the results all occur in different views. I believe that we

can simplify this process by keeping this three-step process within a single view. The app also shows the different percentages for the different estimates it has in the results page. I believe, this is a useful feature as it can show the user alternative guesses if the classifier makes an inaccurate or unsure suggestion. To summarise, I want to develop a system that has a simple classification process and provides adequate result information to the front end such as the probability percentages.

Different types of implementations were also discussed within the review such as cloud based or on-device. The PictureThis app used a cloud-based implementation that requires network connectivity in order to function. The obvious downside to this is that it won't work offline, and network latency will have to be taken into account when the classification process occurs. The flower classification app aims to be fully functional offline by having the model stored on device. Fortunately, because of the advancements of mobile phone hardware, model sizes that may be over 100MB in size are acceptable, which means more accurate models can be stored.

4.1.2 Google Developer Practices

TensorFlow

We will follow the guidelines to transfer a saved model, trained using TensorFlow, into a TensorFlow Lite model, designed for usage on a mobile phone. The Inception V3 model that has been trained and tested in section one is a floating-point model, which means that GPU acceleration can be used in order to decrease inference time (TensorFlow, 2022b). The documentation also provides additional advice on tuning the interpreter to get the best performance. Android Studio, the main IDE used by Android App Developers has tools that can help analyse the real-world performance of the interpreter, which will be key during the development process. By following the advice given by the documentation we can identify key areas of optimisation within the app that should improve performance. Areas that will be considered are performance on the GPU, performance using the CPU with varying number of threads and performance after applying model size reduction techniques.

Android Developers

Apart from general UI interactions like button presses and UI updates, the only additional library support we need to consider from the Android platform is the Camera API. The camera plays a big part in the classification process as it provides the input image information for the model. The application will make use of the newest CameraX API which is designed to be easy-to-use and have support for majority of Android devices including legacy devices down to Android 5 (AndroidDevelopers, 2022). Using this API, we can control the input image dimensions and provide image previews to the user.

Material You

This is more of a secondary objective but conforming to the latest Google UI standards may prove beneficial in creating an easy-to-use UI. Material You is the latest iteration of Android's design language with emphasis on helping "make technology simple and beautiful for everyone" (Material, 2021). Since the app revolves around capturing flowers, widely regarded as beautiful objects, having a simple and pleasing UI seems appropriate.

4.1.3 Functional Requirements

The requirements table can be seen in Appendix A.3.1. It is worth elaborating on them to get a clearer picture of how the final product will function.

Firstly, the app must run on Android devices of minimum API level 26 which is Android version 8.0 and consists of approximately 82.7% of devices (Appendix A.3.3). There are possible ways to make an app that works for iPhone as well using Google's Flutter development kit that allows for deployment to both ecosystems and compatibility with TensorFlow. However, due to the inability to access an iPhone to test the app, it is best to proceed with an Android only implementation.

Through the use of the CameraX API, the app must be able to capture and display images. The camera viewfinder will be visible to the user. When a picture is captured, the view will transition to showing the image that was captured so that the user can clearly see the quality of the image they have captured and whether the subject of the image is in full view.

The app must integrate the Inception V3 model trained in section one. This model will allow for the classification of images captured in the app. The output information which includes the classification results, and their percentage probabilities will be parsed and outputted to the front end.

Results produced by the app must be shown in an acceptable time frame after capturing the image. It is not ideal for the user to have to wait for a long period of time to get the results. I have defined an inference time of less than half a second as acceptable. This falls in line with the estimated times defined on the TensorFlow site which can also be seen in Appendix A.3.4. Within this time period, the app must also display the top three predictions to the front end. If the classifier is unsure about label of the input image, the user can at least see what other possible labels it could be. The predictions will be shown in a clear order from best to third-best prediction. A small thumbnail image of the flower should also accompany the label in the UI so that the user can see if the flower they have captured matches up with the prediction, in case they're not sure of whether the predicted label is correct.

To track if the classifier is recognising images in an acceptable time, a time will be shown within the UI that outputs how long it took from the picture being taken to the image being classified.

4.1.4 Non-functional Requirements

A table of non-functional requirements are also defined in Appendix A.3.2.

The app should function with minimal bugs so that the overall experience of using the app is not hindered by unintended behaviour or crashes. It will also make sure testing of the app's performance goes smoother. If there are crashes or bugs it will mean the testing results are less reliable.

The app will be designed around the Material You design specification which involves producing simple and intuitive user interfaces that conform to the latest Google standards. The aim is to make the app simple to use and not involve too much background UI processes so that the classification process is not hindered in any way.

4.2 Design and Implementation

This section will consist of the design and implementation process of developing the flower classifier app. The objective is to follow the requirements outlined in section 4.1 to produce the app, but first a plan must be outlined to ensure development goes smoothly.

Figure 4.1 shows the simple actions that the user can take in order to capture a flower photo and get results.




Figure 4.1: Flowchart that shows the actions that need to take place to classify a flower in the app.

4.2.1 Classifier

The Inception V3 model has a set of requirements that need to be filled in order to function. This is where using the TensorFlow and CameraX guidelines is important as we have to ensure these pieces work together. The classifier expects an object named `TensorImage`. This object requires the loading of a bitmap image in order to be valid. This means that we must carry out some pre-processing to convert an image into a bitmap, then load the bitmap into a `TensorImage` and then finally process the image to be resized to have a height and width of 299 by 299. The image may also need to be rotated as we get the raw image from the camera sensor, which are not normally in the orientation a user would expect. Figure 4.2 outlines the full process that needs to be carried out to classify a raw image.



Figure 4.2: Diagram that shows the processing the image has to go through until it gets fed to the classifier.

4.2.2 Implementation Process

The environment is the same as what was described in section 3.2.4. The key differences are that the Android Studio IDE will be used for the development of the app. Testing will take place directly on a Samsung Galaxy S20 5G. This device is one of the flagship Samsung phones released in the year 2020 and has decent specifications for a smartphone. The exact specifications can be seen in appendix A.5.1. Android Studio has an in-built tool that can easily process TensorFlow Lite files that have the extension ".tflite". By using the tool, the relevant

app directories are created for the developer, as well as the relevant library dependencies to use TensorFlow.

4.2.3 User Interface Design

While this is not relevant to the base functionality of the application. It is still important to plan a user interface that will allow the user to carry out the simple operation of classifying flowers on their device. Therefore, some rough sketches were produced to plan out the single view that will encompass everything the camera and results interface as seen in figure 4.3. There is also an additional view as seen in figure A.5 that outlines a description page for the flowers.




Figure 4.3: Sketch of the home page.

The home page's main purpose is to capture the flower image and provide the results to the user. The sketch provides that simple outline to follow when designing the interface within the IDE. The viewfinder box at the top should take no more than half the screen to leave enough room for the classification results. There are three rows for the top three results, however the top row is slightly bigger to highlight the top result more than the others. Images for each result are included to help the user see what the predicted flowers actually look like.

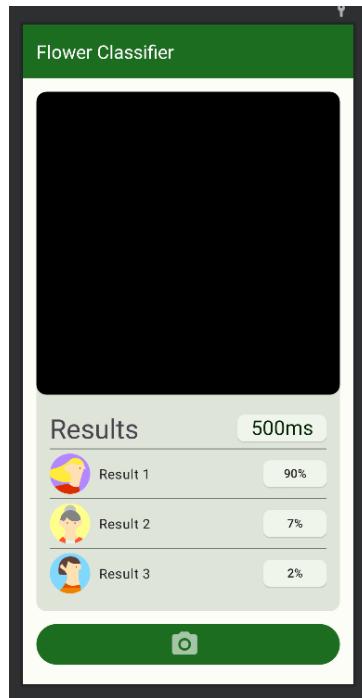


Figure 4.4: IDE preview of the home page.

Figure 4.4 shows the final version of the design that was made within the IDE. It follows the general layout given by the sketch. There are some minor changes with the sizing of some elements to make them stand out more like the camera button. Also, the result elements are all uniform in size to make the image icons clearer. Overall, the design consists of all Material You elements which were located and sized based on the guidelines given by Google.

4.3 System Testing

There are two main aims for the testing portion of this section. The first aim is to ensure the base functionality is working as intended through functionality testing. We will demonstrate real world examples of the app identifying various flowers. Once, the app is established to be working as intended, the next step will be to carry out performance testing to see how the classifier is performing within the app. By using performance profiling, we can focus on specific work threads where the classifier is carrying out the identification process. We will alter certain factors like whether to use the GPU or CPU as well as the number of threads to dedicate to the processor, then discuss the results. Before testing takes place, the phone will use the in-built system clean up tool that clears any unnecessary background apps and processes.

4.3.1 Functionality Testing

This will be split into two formats. The first is to try and identify as many different flowers that are available and report the probabilities for the top three labels as well as the inference time when the app is working on the GPU. The second format is to analyse how the app performs using the same metrics but with the same flower, shot at different angles.

Appendix A.5.2 contains the results of identifying eight different flower species found in the vicinity. All of them bar one has a confirmed species. There are also screenshots that go with

each result within the appendix where one can view the exact angles and lighting conditions the photos were taken in as well as what the other predictions are for that flower. The results show that six of the species have been correctly identified by the classifier. There are five species within that selection that have probabilities of 88 percent and above, each of these flowers are very common in the UK and have distinctive shapes and colours. This suggests the classifier performs very well with these unique flowers, most likely as they don't have too much of a shape and colour crossover with other species. The classifier also works well in identifying the two different dandelion types. These flowers are very different in appearance but are still both dandelions. In the case of the flower identified as a columbine, it is unconfirmed whether the subject flower is in fact a columbine, but by looking at some of the example flowers in the dataset, it is clear why the classifier came to that conclusion. The orchid that is incorrectly identified as a cyclamen has a secondary guess of sweat pea that is also 47%. This suggests that the classifier cannot make a definitive decision of what the flower is, likely due to a lack of orchids within the dataset. Moon orchids are a class within the dataset and are not too dissimilar in appearance to this orchid (PlantToolbox, n.d.). Therefore, it is unclear as to why that class was not suggested over the other classes. Subsequent classification attempts on this same plant yielded consistent high probabilities for the Cyclamen class, suggesting it didn't matter what angle the picture was taken in. The inference times seem consistent using the GPU. It will be interesting to see what the performance will be like when we adjust the classifier settings in the performance testing.

Some additional testing was carried out on the same rose to see how the classifier dealt with different angles and lighting conditions. The first set of tests will include classifying with different camera distances and angles with estimated measurements. The first was the distance test where the multiple shots of the same rose were taken. The camera was angled to have the centre of the rose in the middle of the viewfinder. The rose was always in the same position and location to keep the lighting consistent. The results from this experiment are presented in appendix A.5.3 below. The distance has an effect on the classification accuracy. The prediction at 4cm has a low probability of 39% but still ranks rose as the most likely. This indicates that there are not enough unique features that the classifier can pick up on because the rose is too close to the sensor. When the rose is in full view at 15cm, the entirety of the flower can be seen including its general shape and petal arrangement. As a result, the probability increases by 39%. At 35cm and above, the rose drops below the top 3 predictions, showing that the classifier struggles once more of the background can be seen and there is less data on the subject.

Angles testing was carried out using arbitrary angles of the rose. Images taken of the rose that did not have the camera angled directly at the centre of the flower still ranked rose as the highest prediction as seen in figure [x] and [x] within the appendix. However, the probability decreased to 19% with a side angle of the flower, where the centre of it cannot be seen (figure [x]).

Next, lighting testing was carried out to see how the classifier performed. The full results can be seen in appendix A.5.4. Lighting was measured using the lighting sensor within the camera sensor. A free Android app named DevCheck (GooglePlay, 2022) has access to all sensor readings on the device and was used to measure the light hitting the sensor in lux.

The low light performance of the classifier is better than expected. Initially, an inaccurate result was predicted as it would be difficult to make out the features of the flower with less light. The reason why it may have resulted in a decent prediction is likely due to the great

low-light performance of the camera sensor. When the sensor detects these conditions, it can alter the sensor to let in more light or use software tricks to boost the image clarity. Another unexpected result were the outdoor lighting conditions. These conditions resulted in inaccurate results for the rose, this may be likely due to the colour of the rose becoming whiter in the outdoors which could make it similar to other flowers other than the rose. Additionally, there appears to be less shadows casted within the structure of the rose, which suggests the classifier is having a difficult time making out the edges of the petals due to the flatter look. However, subsequent tests that changed the angle of the flower relative to the direction of sunlight to produce more shadows, still resulted in inaccurate results.

4.3.2 Performance Testing

In this section, the different hardware settings will be tested to see which component is best to carry out the calculations within each layer of the model. The profiler tool within Android Studio will be used to get a detailed report of what is going on, internally. As the profiler pulls in information about all processes going on within the phone, the recording feature will be used to capture specific sections of the profiler output. During the recording, the capture button will be pressed 5 times in a row, the tool will then automatically calculate the average, max, min and standard deviation of the event in milliseconds.

Type	Avg.	Min	Max	SD
CPU (1T)	86.3	79.0	106.8	10.4
CPU (2T)	60.4	50.7	78.1	11.5
CPU (4T)	58.6	44.0	87.2	16.0
CPU (8T)	211.2	191.8	241.2	17.2
GPU	313.1	269.4	313.1	26.8

Table 4.1: Results from inference timings for each hardware delegate configuration.

The full results can be seen in appendix A.5.5 which has the timings of the individual events. Table 4.1 summarises each test with just the calculated statistics from the profiling tool. It shows that using the CPU with four threads has the shortest inference time out of all the test runs. It is surprising to see that the GPU is out-performed by the CPU especially since the GPU can carry out parallelise workloads better than the CPU and should theoretically execute the layer operations quicker. It may still be worth using the GPU over the CPU simply because of the added benefits of accuracy in doing floating point calculations. GPUs are also more energy efficient as they can carry out the same tasks as the CPU, while using less power. Using eight threads does not seem ideal as the average time shoots up. This may be because of scheduling overhead becoming too large due to the lack of CPU resources available. Standard deviation also gets larger as more threads are used, indicating that the inference timings become less consistent, likely due to the state of the CPU at that current time as it juggles other tasks. It may be best to keep using the GPU simply because it's more efficient compared to a CPU. Realistically, a 200ms difference is minuscule in this case. However, it is still worth it for a developer to investigate the best hardware delegate in order to see which one performs best for their problem. For example, if multiple classifications need to occur in sequence, using the GPU may not be ideal as that 200ms difference will slowly add up, especially on slower devices.

4.4 Summary

Overall, the requirements for this app were quite simple. The end product was designed to be a simple utility app that could easily fit as part of the latest Android ecosystem. The requirements have all been sourced from the background research conducted as well as my own input as a software developer. A more extensive set of requirements could be produced by using input from potential users of the app such as gardeners and hobbyists, however, the main aim is to demonstrate that this method of using TensorFlow for deep learning is a very valid approach as a backend for your app. Most of the effort was dedicated into investigating the performance of TensorFlow when put up against the challenging task of classifying flowers. The design and implementation process for the app made sure the scope of the app was kept simple so that the focus could be on testing the performance of the Inception V3 model. Designing complex apps can lead to added performance overhead as more resources are dedicated to tasks that are not related to the classifier. The classifier showed great performance in identifying different flower species in the real world which means that the app functions as it should. Furthermore, while testing, there were no issues with bugs, crashes or performance. It was clear there were some areas that the classifier struggled in when it came to the different possible angles, distances and lighting conditions. This may be due to the lack of variety in the dataset for the rose class. The classifier performed the best when the rose was in clear view under normal lighting conditions. Perhaps, doing some additional random image pre-processing before training could artificially produce more variety, making the classifier better at making predictions under different conditions. Different hardware delegations were also looked at to see what can be configured on the application side to improve the classifier performance. It appears as though using the CPU is faster than the GPU on a smartphone. However, the GPU may still be the best choice for the developer as they can put less stress on the CPU which can carry out other tasks within the app. By purely using the CPU for classification, other tasks within the app may slowdown, making the overall experience for the user worse.

Chapter 5

Conclusion

5.1 Contributions

This project has provided an in-depth insight into the stages of developing a smartphone application that makes use of the Deep Learning technology available. It serves as a good starting point for developers by first providing the relevant background information about the domain. Then, identifying key sources of information that will prove useful to them. Finally, going through the development work required to implement a solution.

The project was able to demonstrate how much of an improvement deep learning is over classical machine learning by comparing the results of two classifiers trained on the same dataset. It was also surprising to see how accessible deep learning is to developers with the tools provided by TensorFlow. Additionally, it's a good opportunity to demonstrate the maturity of classical machine learning resources and code libraries.

In terms of mobile app development, the project shows an updated process of how to train a model through transfer learning and then make it compatible for mobile phones through TensorFlow lite. The steps required to make image classification in general work were also highlighted by demonstrating how to convert a raw image into the format required for the classifier. These steps were generated through studying the available documentation and making use of the CameraX and TensorFlow API to convert the image correctly. There are example apps provided by Google for image classification but contain deprecated camera APIs (see A.26). This solution consists of APIs that are currently supported.

5.2 Critical Appraisal

Due to limitations in the available hardware, the scope of the experiment had to decreased. This means there was limited range of hyperparameters that could be tested for the Inception V3 model as my hardware configuration isn't suited for these tasks and therefore cannot carry out the training processes in a timely manner. Given better hardware resources such as dedicated Nvidia GPUs and larger memory capacities, it would've been possible to conduct more efficient hyper parameter tuning that could lead to better model performance.

It is possible to spend more time focusing on improving the individual models, specifically the SVM model. One could develop a better optimised SVM and carry out additional actions

to improve the image pre-processing pipeline or extract better features. As the aim of the investigation was to compare the approaches from the perspective of a software developer, rather than develop the best possible models, this isn't something that was investigated in-depth. However, there is certainly value in understanding SVMs and other classical machine learning methods more by dedicating more time in implementing them. If one wants to conduct a fairer comparison of DL vs classical ML, they would dedicate more time developing both approaches to get the best possible results for each.

For the development of the app, having user input and testing of potential users would be very helpful in designing a more useful product. There was a lot of focus on implementing an app that uses the DL model effectively rather than the actual usability of the app. As a result, the feature to show additional information about flowers within the app was left out. It would've been useful for testing to show a more featureful app experience to collect results on performance when there are other tasks taking up system resources as well. Then, comparisons could've been made with the results gathered in this project.

5.3 Future Work

There is a lot of potential for deep learning within the smartphone space. There are two areas that we will see improvements in that would require future re-evaluations. The first area are the models such as Inception V4. Researchers will consistently develop new models which can then be potentially re-trained. These models may get larger in size, but could also be more accurate. When new models are produced, future work can be done to see how they compare with the implementation produced in this project. The second area is the dataset. More testing can be done on larger flower datasets that has support for more species in potentially different countries. One could potentially expand on the Oxford Flowers dataset by adding more images into each category.

On-device training is possible through TensorFlow. It would be interesting to see how that compares with the training approach presented in this project. This would effectively remove the need of having a high-powered computer to do the training. There are certainly advantages to this, however, it would be useful to see what the potential caveats with this are and whether it would be a better solution for classifying flowers on a smartphone.

Developers would also be interested in deployment mechanisms of their model to their app. Technically the model is just a file within applications project directory. As long as the input conditions remain the same, the developer could potentially just redeploy an improved model instead of updating the entire application. A potential improvement for the app would be to design such a system so that users could get improved models without the overhead of retraining them, as that has already been carried out by the developer themselves.

5.4 Reflection

This project was a good learning opportunity to explore the capabilities of deep learning on a smartphone. It was surprising to see how effective it is in tackling challenging classification problems such as identifying flowers. By going through the entire process of researching to implementing, I have gained a much better understanding of how to design software of this nature. I believe that the available tools make it surprisingly simple for developers to implement

their own models for their own use cases. By leveraging this technology, overall smartphone experiences will improve over time.

Bibliography

- Agarwal, D., 2021. *Guide for feature extraction techniques* [Online]. Available from: <https://www.analyticsvidhya.com/blog/2021/04/guide-for-feature-extraction-techniques/> [Accessed 2022-01-27].
- AndroidDevelopers, 2022. *Camerax* [Online]. Available from: <https://developer.android.com/training/camerax> [Accessed 2022-04-22].
- Anon, n.d.a. *Convolutional neural networks* [Online]. Available from: https://d2l.ai/chapter_convolutional-neural-networks/lenet.html [Accessed 2021-12-02].
- Anon, n.d.b. *History of the perceptron* [Online]. Available from: <https://home.csulb.edu/~cwallis/artificialn/History.htm> [Accessed 2021-11-30].
- Brownlee, J., 2017. *A gentle introduction to the bag-of-words model* [Online]. Available from: <https://machinelearningmastery.com/gentle-introduction-bag-words-model/> [Accessed 2021-11-06].
- Brownlee, J., 2020a. *4 types of classification tasks in machine learning* [Online]. Available from: <https://machinelearningmastery.com/types-of-classification-in-machine-learning/> [Accessed 2022-01-27].
- Brownlee, J., 2020b. *What is a confusion matrix in machine learning* [Online]. Available from: <https://machinelearningmastery.com/confusion-matrix-machine-learning/#:~:text=A%20confusion%20matrix%20is%20a%20summary%20of%20prediction%20results%20on,key%20to%20the%20confusion%20matrix> [Accessed 2022-03-09].
- Camastra, F. and Vinciarelli, A., n.d. *Machine learning for audio, image and video analysis: Theory and applications*, Advanced Information and Knowledge Processing. 2nd ed. London: Springer London.
- Chapelle, O., Haffner, P. and Vapnik, V.N., 1999. Support vector machines for histogram-based image classification. *IEEE transactions on neural networks*, 10(5), pp.1055–1064.
- Chockwanich, N. and Visoottiviseth, V., 2019. Intrusion detection by deep learning with tensorflow [Online]. *2019 21st international conference on advanced communication technology (icact)*. pp.654–659. Available from: <https://doi.org/10.23919/ICACT.2019.8701969>.
- DataFlair, n.d. *Iris flower classification project using machine learning* [Online]. Available from: <https://data-flair.training/blogs/iris-flower-classification/#:~:text=Iris%20flower%20classification%20is%20a,%2C%20'Petal%20width> [Accessed 2022-01-28].

- Dietterich, T., 1995. Overfitting and undercomputing in machine learning. *Acm comput. surv.* [Online], 27(3), pp.326–327. Available from: <https://doi.org/10.1145/212094.212114>.
- Garcin, C., Joly, A., Bonnet, P., Affouard, A., Lombardo, J.C., Chouet, M., Servajean, M., Lorieul, T. and Salmon, J., 2021. *Pl@ntnet-300k image dataset* (v.1.1) [Online]. Zenodo. Available from: <https://doi.org/10.5281/zenodo.5645731>.
- GeeksForGeeks, 2020. *Choose optimal number of epochs to train a neural network in keras* [Online]. Available from: <https://www.geeksforgeeks.org/choose-optimal-number-of-epochs-to-train-a-neural-network-in-keras/#:~:text=Therefore%20the%20optimal%20number%20of,values%20against%20number%20of%20epochs> [Accessed 2022-04-18].
- Goodfellow, I., Bengio, Y. and Courville, A., 2016. *Deep learning*. MIT press.
- Google, 2021a. *Performance measurement* [Online]. Available from: <https://www.tensorflow.org/lite/performance/measurement> [Accessed 2021-12-02].
- Google, 2021b. *Tensorflow lite* [Online]. Available from: <https://www.tensorflow.org/lite> [Accessed 2021-11-05].
- Google, 2022. *Plantnet plant identification* [Online]. Available from: https://play.google.com/store/apps/details?id=org.plantnet&hl=en_GB&gl=US [Accessed 2022-03-07].
- GoogleCloud, 2022. *Advanced guide to inception v3* [Online]. Available from: <https://cloud.google.com/tpu/docs/inception-v3-advanced> [Accessed 2022-04-21].
- GoogleDevelopers, 2020a. *Classification: Precision and recall* [Online]. Available from: <https://developers.google.com/machine-learning/crash-course/classification/precision-and-recall> [Accessed 2022-04-27].
- GoogleDevelopers, 2020b. *Classification: Roc curve and auc* [Online]. Available from: <https://developers.google.com/machine-learning/crash-course/classification/roc-and-auc> [Accessed 2022-04-16].
- GoogleDevelopers, 2020c. *Descending into ml: Training and loss* [Online]. Available from: <https://developers.google.com/machine-learning/crash-course/descending-into-ml/training-and-loss#:~:text=Loss%20is%20the%20penalty%20for,otherwise%20the%20loss%20is%20greater> [Accessed 2022-04-16].
- GoogleDevelopers, 2020d. *Multi-class neural networks: Softmax* [Online]. Available from: <https://developers.google.com/machine-learning/crash-course/multi-class-neural-networks/softmax> [Accessed 2022-04-21].
- GooglePlay, 2022. *Devcheck* [Online]. Available from: https://play.google.com/store/apps/details?id=flar2.devcheck&hl=en_GB&gl=US [Accessed 2022-04-26].
- GSMArena, 2021. *Gsm arena* [Online]. Available from: <https://www.gsmarena.com/> [Accessed 2021-11-30].
- Gupta, M., 2021. *Google tensor is a milestone for machine learning* [Online]. Available from: <https://blog.google/products/pixel/introducing-google-tensor/> [Accessed 2021-11-30].

- Joly, A., Bonnet, P., Goëau, H., Barbe, J., Selmi, S., Champ, J., Dufour-Kowalski, S., Affouard, A., Carré, J., Molino, J.F., Boujema, N. and Barthélémy, D., 2016. A look inside the PI@ntNet experience. *Multimedia Systems* [Online], 22(6), pp.751–766. Available from: <https://doi.org/10.1007/s00530-015-0462-9>.
- Kavlakoglu, E., 2020. *Ai vs. machine learning vs. deep learning vs. neural networks: What's the difference?* [Online]. Available from: <https://www.ibm.com/cloud/blog/ai-vs-machine-learning-vs-deep-learning-vs-neural-networks> [Accessed 2021-12-02].
- Keras, n.d.a. *About keras* [Online]. Available from: <https://keras.io/about/#:~:text=Keras%20is%20the%20high%2Dlevel,solutions%20with%20high%20iteration%20velocity> [Accessed 2022-04-20].
- Keras, n.d.b. *Concatenate layer* [Online]. Available from: https://keras.io/api/layers/merging_layers/concatenate/ [Accessed 2022-04-21].
- Korstanje, J., 2021. *The f1 score* [Online]. Available from: <https://towardsdatascience.com/the-f1-score-bec2bbc38aa6> [Accessed 2021-12-02].
- Kulendran, M., Lim, M., Laws, G., Chow, A., Nehme, J., Darzi, A. and Purkayastha, S., 2014. Surgical smartphone applications across different platforms: Their evolution, uses, and users. *Surgical innovation*, 21(4), pp.427–440.
- LeCun, Y., Bengio, Y. and Hinton, G., 2015. Deep learning. *nature*, 521(7553), pp.436–444.
- Liu, Z., Lin, Y. and Sun, M., 2020. *Representation learning for natural language processing*. Springer Nature.
- Maithani, M., 2021. *Guide to tensorflow keras optimizers* [Online]. Available from: <https://analyticsindiamag.com/guide-to-tensorflow-keras-optimizers/> [Accessed 2022-04-18].
- Material, 2021. *Unveiling material you* [Online]. Available from: <https://material.io/blog/announcing-material-you> [Accessed 2022-04-23].
- MathWorks, n.d. *Texture analysis* [Online]. Available from: <https://www.mathworks.com/help/images/texture-analysis-1.html> [Accessed 2022-03-09].
- Mete, B.R. and Ensari, T., 2019. Flower classification with deep cnn and machine learning algorithms [Online]. *2019 3rd international symposium on multidisciplinary studies and innovative technologies (ismsit)*. pp.1–5. Available from: <https://doi.org/10.1109/ISMSIT.2019.8932908>.
- Mingqiang, Y., Kidiyo, K., Joseph, R. et al., 2008. A survey of shape feature extraction techniques. *Pattern recognition*, 15(7), pp.43–90.
- Mohammed, M., 2017. *Machine learning : algorithms and applications*. 1st ed. Boca Raton: CRC Press.
- Mohan, S., 2020. *Image classification using bag of visual words model* [Online]. Available from: <https://machinelearningknowledge.ai/image-classification-using-bag-of-visual-words-model/> [Accessed 2022-04-21].

- Nilsback, M.E. and Zisserman, A., 2008. Automated flower classification over a large number of classes [Online]. *2008 sixth indian conference on computer vision, graphics image processing*. pp.722–729. Available from: <https://doi.org/10.1109/ICVGIP.2008.47>.
- Noble, W.S., 2006. What is a support vector machine? *Nature biotechnology*, 24(12), pp.1565–1567.
- O'Dea, S., 2021. *Global smartphone market share from 4th quarter 2009 to 3rd quarter 2021* [Online]. Available from: <https://www.statista.com/statistics/271496/global-market-share-held-by-smartphone-vendors-since-4th-quarter-2009/> [Accessed 2021-11-30].
- Olascoaga, L.I.G., Meetr, W. and Verhelst, M., 2021. *Hardware-aware probabilistic machine learning models: Learning, inference and use cases*. Springer.
- PapersWithCode, 2017. *inaturalist* [Online]. Available from: <https://paperswithcode.com/dataset/inaturalist> [Accessed 2022-04-14].
- Perlato, A., n.d. *The learning rate* [Online]. Available from: <https://www.andreaperlato.com/theorypost/the-learning-rate/> [Accessed 2022-04-18].
- PictureThisAI, n.d. *Picturethis ai* [Online]. Available from: <https://www.picturethisai.com/> [Accessed 2022-03-07].
- PlantToolbox, n.d. *Phalaenopsis* [Online]. Available from: <https://plants.ces.ncsu.edu/plants/phalaenopsis/> [Accessed 2022-04-25].
- Pratap Singh, S., n.d. *Fully connected layer: The brute force layer of a machine learning model* [Online]. Available from: <https://iq.opengenus.org/fully-connected-layer/> [Accessed 2022-04-21].
- Rosebrook, A., 2016. *k-nn classifier for image classification* [Online]. Available from: <https://www.pyimagesearch.com/2016/08/08/k-nn-classifier-for-image-classification/> [Accessed 2021-12-02].
- Rosenblatt, F., 1961. *Principles of neurodynamics. perceptrons and the theory of brain mechanisms*. Cornell Aeronautical Lab Inc Buffalo NY.
- Saha, S., 2018. *A comprehensive guide to convolutional neural networks — the eli5 way* [Online]. Available from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53> [Accessed 2021-12-02].
- Sayak, P., 2018. *Hyperparameter optimization in machine learning models* [Online]. Available from: <https://www.datacamp.com/community/tutorials/parameter-optimization-machine-learning-models> [Accessed 2022-04-20].
- Scarpino, M., 2018. *Tensorflow for dummies*, For dummies. Newark: Wiley.
- Scikit, n.d. *Precision score* [Online]. Available from: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.precision_score.html [Accessed 2022-04-17].
- Sebastian, 2021. *Dropout regularization in neural networks: How it works and when to use it* [Online]. Available from: <https://programmatically.com/>

- dropout-regularization-in-neural-networks-how-it-works-and-when-to-use-it/ [Accessed 2022-04-18].
- Seeger, M., 2000. Learning with labeled and unlabeled data. *n/a*.
- Shung, K.P., 2018. *Accuracy, precision, recall or f1?* [Online]. Available from: <https://towardsdatascience.com/accuracy-precision-recall-or-f1-331fb37c5cb9> [Accessed 2021-12-02].
- Szeliski, R., 2011. *Computer vision : Algorithms and applications*, Texts in Computer Science. 1st ed. London: Springer London : Imprint: Springer.
- TeaPearce, 2021. *oxford flowers102 bad splits* [Online]. Available from: <https://github.com/tensorflow/datasets/issues/3022> [Accessed 2022-04-17].
- TensorFlow, 2021. *Hosted models* [Online]. Available from: https://www.tensorflow.org/lite/guide/hosted_models [Accessed 2022-04-22].
- TensorFlow, 2022a. *oxford flowers102* [Online]. Available from: https://www.tensorflow.org/datasets/catalog/oxford_flowers102 [Accessed 2022-04-17].
- TensorFlow, 2022b. *Tensorflow lite on gpu* [Online]. Available from: https://www.tensorflow.org/lite/performance/gpu_advanced [Accessed 2022-04-26].
- Three, 2021. *Samsung galaxy s21 ultra 5g* [Online]. Available from: <http://www.three.co.uk/samsung/galaxy-s21-ultra-5g?colour=phantom%20black&memory=128&paym=true> [Accessed 2021-12-02].
- Tian, D., 2013. A review on image feature extraction and representation techniques. *International journal of multimedia and ubiquitous engineering*, 8, pp.385–395.
- Triggs, R. and Simons, H., 2021. *Google tensor vs snapdragon 888 series: How the pixel 6's chip shapes up* [Online]. Available from: <https://www.androidauthority.com/google-tensor-vs-snapdragon-888-3025332/> [Accessed 2021-11-30].
- Turk, M.A. and Pentland, A.P., 1991. Face recognition using eigenfaces. *Proceedings. 1991 ieee computer society conference on computer vision and pattern recognition*. IEEE Computer Society, pp.586–587.
- Williams, N., Zander, S. and Armitage, G., 2006. A preliminary performance comparison of five machine learning algorithms for practical ip traffic flow classification. *Sigcomm comput. commun. rev.* [Online], 36(5), p.5–16. Available from: <https://doi.org/10.1145/1163593.1163596>.
- Xia, X., Xu, C. and Nan, B., 2017. Inception-v3 for flower classification [Online]. *2017 2nd international conference on image, vision and computing (icivc)*. pp.783–787. Available from: <https://doi.org/10.1109/ICIVC.2017.7984661>.
- Xin, Y., Kong, L., Liu, Z., Chen, Y., Li, Y., Zhu, H., Gao, M., Hou, H. and Wang, C., 2018. Machine learning and deep learning methods for cybersecurity. *ieee access* [Online], 6, pp.35365–35381. Available from: <https://doi.org/10.1109/ACCESS.2018.2836950>.
- Zhang, H., 2004. The optimality of naive bayes. *Aa*, 1(2), p.3.
- Zhu, X.J., 2005. Semi-supervised learning literature survey. *n/a*.

Appendix A

Appendix

A.1 Table of Smartphones

Phone (Year)	Processor	Storage (GB)	Memory (GB)	Cameras (MP)
S2 (2011)	Dual-core 1.2 GHz	32	1	8
S3 (2012)	Quad-core 1.4 GHz	64	1	8
S4 (2013)	Octa-Core (4x1.6 GHz, 4x1.2 GHz)	64	2	13
S5 (2014)	Quad-Core 2.5 GHz	32	2	16
S6 edge+ (2015)	Octa-Core (4x2.1 GHz, 4x 1.5 GHz)	64	4	16
S7 edge (2016)	Octa-Core (4x 2.3 GHz, 4x 1.6 GHz)	128	4	12
S8+ (2017)	Octa-Core (4x 2.35 GHz, 4x 1.9 GHz)	128	6	12
S9+ (2018)	Octa-Core (4x 2.8 GHz, 4x 1.7 GHz)	256	6	12/12
S10+ (2019)	Octa-Core (4x 2.84 GHz, 4x 1.78 GHz)	1024	12	12/12/16
S20 Ultra 5G (2020)	Octa-Core (1x 2.84 GHz, 3x 2.42 GHz, 4x 1.8 GHz)	512	16	0.3/12/48/108
S21 Ultra 5G (2021)	Octa-Core (1x 2.84 GHz, 3x 2.42 GHz, 4x 1.8 GHz)	512	16	10/10/12/108

Table A.1: All data is sourced from GSMArena (2021).

A.2 Inception V3 Results Graphs

A.2.1 Full HP Results

Trial ID	Show Metrics	dropout	learning rate	Accuracy
63ae00604944d...	<input type="checkbox"/>	0.20000	0.0010000	0.91863
fc373df789b3f6...	<input type="checkbox"/>	0.30000	0.0010000	0.91863
38d32e552556f...	<input type="checkbox"/>	0.40000	0.0010000	0.91569
24b5edbac9ee0...	<input type="checkbox"/>	0.40000	0.010000	0.90882
a4514be3c6213...	<input type="checkbox"/>	0.30000	0.010000	0.90196
a55c7d9cd3858...	<input type="checkbox"/>	0.20000	0.010000	0.89314
596bae5bda236...	<input type="checkbox"/>	0.20000	0.00010000	0.33725
89364737c82ec...	<input type="checkbox"/>	0.40000	0.00010000	0.30196
1aac363301725...	<input type="checkbox"/>	0.30000	0.00010000	0.29510

Figure A.1: Table of full HP training results as seen in TensorBoard.

A.2.2 Loss




Figure A.2: Graph of loss against the number of training steps.

A.3 Requirements

Sources:

- SS: Saahil Shihaz
- AD: Android Developers
- TF: TensorFlow
- MY: Material You

Priorities:

- H: High
- M: Medium
- L: Low

A.3.1 Functional Requirements

No.	Description	Source	Priority
1	The app must run on an Android device.	SS	H
2	The app must be able to capture an image.	SS, AD	H
2.1	The app must be able to display a captured image.	SS, AD	H
3	The app must use the Inception V3 model trained in section one to process captured images.	SS, TF	H
4	The app must display the guesses produced by the model within 500 milliseconds.	SS	H
4.1	The app must show the top three guesses as well as their individual percentage probabilities ranked from highest to lowest.	SS	H
4.2	The app should show a small preview of the flowers next to the guesses.	SS	M
5	The app must display the time it takes to process an image in milliseconds.	SS	H
6	The app could contain additional pages that provide more pictures and information about the flower that can be accessed by clicking on the flower's icon.	SS	L

A.3.2 Non-functional Requirements

No.	Description	Source	Priority
1	The app must function with minimal bugs.	SS	H
2	The app must follow the Material You design specifications.	MY	L

A.3.3 SDK percentage




Figure A.3: The percentage of devices that API level 8 will run according to Android Studio.

A.3.4 Model performance list

Model name	Paper and model	Model size	Top-1 accuracy	Top-5 accuracy	CPU, 4 threads	GPU	NNAPI
DenseNet	paper , tflite&pb	43.6 Mb	64.2%	85.6%	195 ms	60 ms	1656 ms
SqueezeNet	paper , tflite&pb	5.0 Mb	49.0%	72.9%	36 ms	9.5 ms	18.5 ms
NASNet mobile	paper , tflite&pb	21.4 Mb	73.9%	91.5%	56 ms	--	102 ms
NASNet large	paper , tflite&pb	355.3 Mb	82.6%	96.1%	1170 ms	--	648 ms
ResNet_V2_101	paper , tflite&pb	178.3 Mb	76.8%	93.6%	526 ms	92 ms	1572 ms
Inception_V3	paper , tflite&pb	95.3 Mb	77.9%	93.8%	249 ms	56 ms	148 ms
Inception_V4	paper , tflite&pb	170.7 Mb	80.1%	95.1%	486 ms	93 ms	291 ms
Inception_ResNet_V2	paper , tflite&pb	121.0 Mb	77.5%	94.0%	422 ms	100 ms	201 ms
Mobilenet_V1_0.25_128	paper , tflite&pb	1.9 Mb	41.4%	66.2%	1.2 ms	1.6 ms	3 ms
Mobilenet_V1_0.25_160	paper , tflite&pb	1.9 Mb	45.4%	70.2%	1.7 ms	1.7 ms	3.2 ms
Mobilenet_V1_0.25_192	paper , tflite&pb	1.9 Mb	47.1%	72.0%	2.4 ms	1.8 ms	3.0 ms
Mobilenet_V1_0.25_224	paper , tflite&pb	1.9 Mb	49.7%	74.1%	3.3 ms	1.8 ms	3.6 ms
Mobilenet_V1_0.50_128	paper , tflite&pb	5.3 Mb	56.2%	79.3%	3.0 ms	1.7 ms	3.2 ms
Mobilenet_V1_0.50_160	paper , tflite&pb	5.3 Mb	59.0%	81.8%	4.4 ms	2.0 ms	4.0 ms
Mobilenet_V1_0.50_192	paper , tflite&pb	5.3 Mb	61.7%	83.5%	6.0 ms	2.5 ms	4.8 ms
Mobilenet_V1_0.50_224	paper , tflite&pb	5.3 Mb	63.2%	84.9%	7.9 ms	2.8 ms	6.1 ms
Mobilenet_V1_0.75_128	paper , tflite&pb	10.3 Mb	62.0%	83.8%	5.5 ms	2.6 ms	5.1 ms
Mobilenet_V1_0.75_160	paper , tflite&pb	10.3 Mb	65.2%	85.9%	8.2 ms	3.1 ms	6.3 ms
Mobilenet_V1_0.75_192	paper , tflite&pb	10.3 Mb	67.1%	87.2%	11.0 ms	4.5 ms	7.2 ms
Mobilenet_V1_0.75_224	paper , tflite&pb	10.3 Mb	68.3%	88.1%	14.6 ms	4.9 ms	9.9 ms
Mobilenet_V1_1.0_128	paper , tflite&pb	16.9 Mb	65.2%	85.7%	9.0 ms	4.4 ms	6.3 ms
Mobilenet_V1_1.0_160	paper , tflite&pb	16.9 Mb	68.0%	87.7%	13.4 ms	5.0 ms	8.4 ms
Mobilenet_V1_1.0_192	paper , tflite&pb	16.9 Mb	69.9%	89.1%	18.1 ms	6.3 ms	10.6 ms
Mobilenet_V1_1.0_224	paper , tflite&pb	16.9 Mb	71.0%	89.9%	24.0 ms	6.5 ms	13.8 ms
Mobilenet_V2_1.0_224	paper , tflite&pb	14.0 Mb	71.8%	90.6%	17.5 ms	6.2 ms	11.23 ms

Figure A.4: The list of models and their performance metrics on the TensorFlow site (TensorFlow, 2021).

A.4 Design

A.4.1 Prototype




Figure A.5: Sketch of the flower information page

A.5 Testing

A.5.1 Testing device specifications

Samsung Galaxy S20 5G

Apr 27, 2022 5:57

Hardware

exynos990

Cores: 8

CPU:

2 x M5

2 x Cortex-A76

4 x Cortex-A55

Process: 7 nm LPP

Frequencies:

442 MHz - 2002 MHz

507 MHz - 2504 MHz

546 MHz - 2730 MHz

GRAPHICS

Vendor: ARM

GPU: Mali-G77

OpenGL: OpenGL ES 3.2

Max frequency: 800 MHz

Resolution: 2400 x 1080

Screen density: 424.48477 ppi

Screen size (estimated): 6.2 in / 158 mm

RAM

RAM size: 12 GB

Type: LPDDR5 2750 MHz

Bandwidth: 44 GB/s

Channels: 16-bit quad channel

Storage

Size: 128 GB

Filesystem: sdcarddfs

DEVICE

Model: SM-G981B

Codename: x1s

Manufacturer: samsung

Manufacturing date: October 5, 2020

System

Android Version: Android 12

Build: SP1A.210812.016.G981BXXSDFVC9

ROM base: G981BXXSDFVC9

Security patch: April 1, 2022

Architecture: aarch64 (64-bit)

Instruction sets: arm64-v8a armeabi-v7a armeabi

Kernel: Linux version 4.19.87-23725627 (dpi@21DJ7D03) (Android (dev based on r349610) clang version 8.0.8 (based on LLVM 8.0.8svn))

Battery

Technology: Li-ion

Health: Good

Capacity (reported by system): 3880 mAh

CAMERA

Resolution: 12.2 MP (4032x3024)

Focal length: 2.2 mm

35mm equivalent focal length: 13.5 mm

Sensor size: 5.64 x 4.23 mm

Crop factor: 6.1x

Field of view: 104.1 degrees
 Pixel size: 1.40 micro metres
 Aperture: 2.2
 Shutter speed: 1/11764 - 1/10 s
 RAW mode: No
 ISO sensitivity range: 50 - 3200
 RAW mode: Supported
 Optical image stabilization: No
 Front camera: 7.1 MP (3216x2208)

A.5.2 Results of functionality testing

Figure	Prediction	Prob. (%)	Actual	Time (ms)
A.6	Anthurium	78	Anthurium	351
A.7	Columbine	48	Unknown	376
A.8	Cyclamen	47	Orchid	364
A.9	Daffodil	93	Daffodil	378
A.10	Oxeye Daisy	96	Oxeye Daisy	379
A.11	Dandelion	100	Dandelion	393
A.12	Dandelion	88	Dandelion	373
A.13	Rose	97	Rose	316

Table A.2: Table of predictions for 8 different flowers found in the vicinity.




Figure A.6: Classifying an Anthurium using the app.




Figure A.7: Classifying a Columbine using the app.




Figure A.8: Classifying an Orchid using the app.




Figure A.9: Classifying a Daffodil using the app.




Figure A.10: Classifying a Daisy using the app.




Figure A.11: Classifying a yellow Dandelion using the app.




Figure A.12: Classifying a white Dandelion using the app.




Figure A.13: Classifying a Rose using the app.

A.5.3 Results of distance testing

Figure	Distance (cm)	Prediction	Prob. (%)
A.14	4	Rose	39
A.15	15	Rose	78
A.16	35	Sword Lily	56
A.17	60	Carnation	72

Table A.3: Results from attempting to identify a rose at different distances.




Figure A.14: Classifying a Rose at 4cm distance using the app.




Figure A.15: Classifying a Rose at 15cm distance using the app.




Figure A.16: Classifying a Rose at 35cm distance using the app.




Figure A.17: Classifying a Rose at 60cm distance using the app.

A.5.4 Results of lighting testing

Figure	Lighting (lux)	Prediction	Prob. (%)
A.18	2.9	Rose	34
A.19	25.4	Rose	68
A.20	5333.6	Thorn Apple	45

Table A.4: Results from attempting to identify a rose at different lighting conditions.




Figure A.18: Classifying a Rose in dark conditions using the app.




Figure A.19: Classifying a Rose in normal conditions using the app.




Figure A.20: Classifying a Rose in outdoor conditions using the app.

A.5.5 Performance Timings

5 Count	86.25 ms Average	106.79 ms Max	79 ms Min	10.35 ms Std Dev
Longest running occurrences (select row to navigate)				
Start Time	Name	Wall Duration		
00:01.048	runInference	106.79 ms		
00:02.649	runInference	82.68 ms		
00:05.461	runInference	82.12 ms		
00:04.036	runInference	80.67 ms		
00:06.871	runInference	79 ms		

Figure A.21: Execution timings of the CPU with 1 Thread.

5 Count	60.35 ms Average	78.08 ms Max	50.69 ms Min	11.46 ms Std Dev
Longest running occurrences (select row to navigate)				
Start Time	Name	Wall Duration		
00:01.371	runInference	78.08 ms		
00:06.807	runInference	69.95 ms		
00:02.727	runInference	52.08 ms		
00:05.331	runInference	50.94 ms		
00:03.991	runInference	50.69 ms		

Figure A.22: Execution timings of the CPU with 2 Threads.

5 Count	58.57 ms Average	87.17 ms Max	43.99 ms Min	15.96 ms Std Dev
Longest running occurrences (select row to navigate)				
Start Time		Name		Wall Duration
00:03.768		runInference		87.17 ms
00:00.926		runInference		64.79 ms
00:05.093		runInference		48.58 ms
00:02.328		runInference		48.31 ms
00:06.596		runInference		43.99 ms

Figure A.23: Execution timings of the CPU with 4 Threads.


5 Count	211.18 ms Average	241.15 ms Max	191.79 ms Min	17.17 ms Std Dev
Longest running occurrences (select row to navigate)				
Start Time		Name		Wall Duration
00:03.113		runInference		241.15 ms
00:01.611		runInference		217.71 ms
00:04.712		runInference		204.9 ms
00:06.149		runInference		200.36 ms
00:07.634		runInference		191.79 ms

Figure A.24: Execution timings of the CPU with 8 Threads.

5 Count	313.12 ms Average	347.58 ms Max	269.38 ms Min	26.79 ms Std Dev
Longest running occurrences (select row to navigate)				
Start Time		Name		Wall Duration
00:01.226		runInference		347.58 ms
00:07.009		runInference		330.99 ms
00:08.982		runInference		316.95 ms
00:03.246		runInference		300.71 ms
00:05.056		runInference		269.38 ms

Figure A.25: Execution timings of the GPU.

A.6 Google Image Classifier Example



The screenshot shows a code editor window with a list of errors and warnings. The title bar indicates the file is `CameraActivity.java`. The list of issues includes:

- Warning: 'android.hardware.Camera' is deprecated as of API 21: Android 5.0 (Lollipop) :62
- Warning: Field 'yuvBytes' may be 'final' :76
- Warning: Field can be converted to a local variable :81
- Warning: Unnecessary; SDK_INT is never < 21 :134
- Warning: 'removeGlobalOnLayoutListener(android.view.ViewTreeObserver.OnGlobalLayoutListener)' is deprecated as of API 16: Android 4.1 (Jelly Bean) :135
- Warning: Switch statement on an 'int' with known associated constant missing case 'BottomSheetBehavior.STATE_HALF_EXPANDED' :151
- Warning: Method 'getLuminanceStride()' is never used :205
- Warning: Method 'getLuminance()' is never used :209
- Warning: 'android.hardware.Camera' is deprecated as of API 21: Android 5.0 (Lollipop) :215
- Warning: 'android.hardware.Camera' is deprecated as of API 21: Android 5.0 (Lollipop) :224
- Warning: 'android.hardware.Camera.Size' is deprecated as of API 21: Android 5.0 (Lollipop) :224
- Warning: 'previewWidth' should probably not be assigned to 'yRowStride' :237
- Warning: Anonymous new Runnable() can be replaced with lambda :240
- Warning: Anonymous new Runnable() can be replaced with lambda :248
- Warning: Anonymous new Runnable() can be replaced with lambda :288
- Warning: Anonymous new Runnable() can be replaced with lambda :305
- Warning: Not annotated parameter overrides @NonNull parameter :374

Figure A.26: Output from the IDE when loading the example project.