

Literature and Technology Review

Saahil Shihaz

December 3, 2021

Contents

1	Introduction	2
2	Mobile Machine Learning	2
2.1	Evolution	3
2.2	Where does this lead us to, today?	4
3	Computer Vision	5
3.1	History	5
4	Machine Learning	6
4.1	Feature Extraction	7
4.2	Classification	7
5	Deep Learning	9
5.1	Neurons and Perceptrons	9
5.2	Convolutional Neural Networks	10
5.3	ML vs DL	12
6	Flower Classification	13
6.1	Existing Methods	13
6.2	Existing Apps	14
7	Evaluation	15
8	Final Thoughts	16

1 Introduction

Ground-breaking achievements in technology and specifically machine learning have given us the tools and capabilities to tackle the key problem of image recognition. What this project aims to demonstrate is the application of deep learning within a mobile application to recognise flowers. It will be interesting to see how deep learning performs on mobile hardware which is generally less powerful than desktop PCs and laptops. Additionally, we will get to explore what sort of optimisations need to take place in order to get a feasible mobile based solution. The limited computing resources that we have to work with when creating our solutions is what makes this project challenging. Mobile devices typically contain smaller mobile processors, limited storage and batteries. These restrictions are in place to make mobile phones more efficient and to ensure that they last longer when not connected to a power source. In addition to this, the subject of machine and deep learning is complex and many find the concepts challenging to understand. What this literature review aims to do is identify key sources of information to help breaking down the underlying subject and discuss the quality of the research available.

2 Mobile Machine Learning

The potential of smartphones has still not been fully realised. With advancements in machine learning we can take the capabilities of smartphones to the next level by leveraging the advanced hardware within them to carry out complex tasks. This section will identify the key milestones within smartphone technology and how it leads us to integrating machine learning to make full use of the hardware.



Figure 1: Samsung Galaxy S21 Ultra 5G: Boasting an impressive array of camera sensors on the back (Three, 2021).

2.1 Evolution

Firstly, we want to look at the last 10 years of technological advancements within the smartphone space. With this information, we can hopefully gain some insight into the how much their capabilities have evolved. To do this, we will collect key aspects of specification data from the Samsung Galaxy flagship line of smartphones. Samsung currently hold the top spot in global market share at 20.8% as of Q3' 2021, this position is typically held by Apple or Samsung and can vary from a quarter to quarter basis (O'Dea, 2021).

Phone (Year)	Processor	Storage (GB)	Memory (GB)	Cameras (MP)
S2 (2011)	Dual-core 1.2 GHz	32	1	8
S3 (2012)	Quad-core 1.4 GHz	64	1	8
S4 (2013)	Octa-Core (4x1.6 GHz, 4x1.2 GHz)	64	2	13
S5 (2014)	Quad-Core 2.5 GHz	32	2	16
S6 edge+ (2015)	Octa-Core (4x2.1 GHz, 4x 1.5 GHz)	64	4	16
S7 edge (2016)	Octa-Core (4x 2.3 GHz, 4x 1.6 GHz)	128	4	12
S8+ (2017)	Octa-Core (4x 2.35 GHz, 4x 1.9 GHz)	128	6	12
S9+ (2018)	Octa-Core (4x 2.8 GHz, 4x 1.7 GHz)	256	6	12/12
S10+ (2019)	Octa-Core (4x 2.84 GHz, 4x 1.78 GHz)	1024	12	12/12/16
S20 Ultra 5G (2020)	Octa-Core (1x 2.84 GHz, 3x 2.42 GHz, 4x 1.8 GHz)	512	16	0.3/12/48/108
S21 Ultra 5G (2021)	Octa-Core (1x 2.84 GHz, 3x 2.42 GHz, 4x 1.8 GHz)	512	16	10/10/12/108

Table 1: All data is sourced from GSMArena, 2021

We can see a clear increase in smartphone capability in multiple categories like the processor speed, core count, storage, memory and camera capabilities. There are of course many more different areas that are not listed like sensors, screen size and battery life which have also seen massive improvements over the last 10 years. We have only been seeing improvements in this space, therefore we can assume that we will continue to see improvements in the near future. What we must also consider is price and accessibility, flagship smartphones represent the top of the line offerings from each smartphone manufacturer and are of course priced as such. Low-end to mid-end smartphones are still more capable than their predecessors albeit their spec sheet may not be as impressive as high-end versions in the same generation. Therefore, we must ensure some level of scalability within our machine and deep learning processes. How can we make sure our processes can run efficiently on lower end hardware as well as high end hardware?

Kulendran et al. (2014), highlights how improvements in smartphones have created a boom in the number of smartphone based applications designed to aid surgeons and patients in multiple facets of the medical industry like plastic, orthopaedic and general surgery. They conduct an expansive review of different solutions and analyse how the evolution of smartphones got them to the point that make them extremely useful as a tool to aid us. Ultimately, what this project aims to do is provide a robust software solution to recognise flower species on a smartphone, but we cannot ignore the fact that smartphones have come a long way in the hardware and operating system space to allow us to even conceive of a system.

2.2 Where does this lead us to, today?

ML and AI has become such an important part of smartphones that manufacturers now have dedicated processors for ML and AI tasks. Google includes a Tensor Processing Unit (TPU) in their Pixel line of phones (Triggs and Simons, 2021). Samsung, Qualcomm and Apple use their own solutions for machine learning processing by having their own bespoke processors. These processors are used to compute specific actions that require the decision making and accuracy capabilities of machine learning. Google Tensor in particular aids tasks such as speech recognition that is accurate but not taxing, therefore saving battery life. Tensor also applies to processing photographs and provides additional features to videos (Gupta, 2021). With such a focus on smartphones, to the point that they get dedicated hardware for ML, we should be seeing a huge increase in applications that integrate ML in some way, as well as the entire process of designing and implementing such solutions being carried out more rapidly, as developers learn to leverage the hardware.

3 Computer Vision

Since we are working with analysing images, the area of computer vision plays a big part in our research. In order to identify flower species we must first discuss techniques to analyse the incoming image data to make predictions using ML and DL.

3.1 History

Szeliski (2011) outlines significant occurrences in each decade starting from the 70s, thought to be the beginning of computer vision, all the way through to the 2000s. In the early 70s, researchers sought to emulate human intelligence in a machine by first solving the visual problem. It was hypothesised that if a computer could first recognize objects in the real world that it could then move onto the next step of using reasoning and problem solving at a high level. The first processes conducted to understand the 3D world were to extract edges to recognize 3D objects from 2D lines in an image.

The 80s were described to have a lot more focus on mathematical techniques for analysing scenes. Various algorithms and models were conceived as well as improvements in the contour and edge detection space. Researchers found that a lot of these algorithms could be thought of as “optimization problems” when they were described using the same mathematical framework.

We see more improvements in the field during the 90s including the production of 3D surfaces, tracking and image segmentation. However, what is probably more relevant to this project is statistical learning techniques that also started to appear during this decade. In 1991, we see a paper by Turk and Pentland (1991) that described the concept of “eigenfaces”. These are the product of converting images of faces into feature images. These feature images are essentially the training set. Recognition occurs “by projecting a new image into the sub-space spanned by the eigenfaces”. The new face is then classified by comparing it’s position relative to the known set of faces. Emphasis was placed on the limiting the scope of the allowed images, as such the system was trained and ready to accept profile straight-on images of the subject. In addition to that, they aimed to have the system compute a result in a reasonable time, which of course, is one of the goals of this project. The research hoped to improve on it’s predecessors that used, at the time, traditional methods of recognising features such as eyes’, noses and mouths and their relative position to each other. The work done with eigenfaces shows great similarities with the machine learning techniques we see today, by essentially creating feature vectors and comparing the distance of known vectors in the same space.

Szeliski (2011) continues with their insight into the 2000s where see the various improvements like more efficient algorithms and what finally dominates the latter half of the 2000s; applying machine learning techniques to computer

vision to aid visual recognition research.

4 Machine Learning

This project will be using ML techniques to compare efficiency and accuracy to the more evolved deep learning. What we must first consider is how machine learning works in the context of computer vision. Camastra and Vinciarelli (2015) summarise this and broke down ML development around three primary research points:

- Task-Oriented Studies, improving performance of learning systems in a predetermined set of tasks.
- Cognitive Simulation, emulating the human brain and designing processes around the human thought process.
- Theoretical Analysis, “the theoretical investigation of possible learning methods and algorithms independently of application domain”.

They also produce a taxonomy to represent the balance of two entities they describe: the “teacher” and the “learner”. The teacher, being the programmer, the one that designs the learning process and the learner being the computer system. The idea of inference is also introduced where a system can derive knowledge from previous observations. The taxonomy breaks down the amount of work that both the “learner” and the “teacher” need to do into four categories: Rote Learning, Learning from instruction, Learning by analogy and Learning from examples.

What we are more interested in is learning from examples where the “learner” infers the most out of the other categories in the taxonomy. The idea of the “learning problem” is introduced where the system needs to find a “general rule that explains the data given only a sample of limited size”. Learning techniques are broken down into four more categories: Supervised learning, Reinforcement learning, Unsupervised learning, Semi-supervised learning.

Zhu (2005) highlights semi-supervised learning in their survey as the combination of supervised and unsupervised learning where we use both labelled and unlabelled data for training of the classifier. They point to the survey done by Seeger (2000) in particular that provides more insight into the concept of semi-supervised learning. Their rationale for the concept in general was applying the ability for a system to make predictions based on knowledge it doesn’t have. A supervised system has all labelled data to aid it’s training therefore it’s basis on making predictions is described as a “security belt” by Seeger. The model will basically make predictions within its limited scope, what we would call “overfitting” (Dietterich, 1995). Unsupervised learning heavily relies on prior assumptions for their final result this is because it doesn’t have a knowledge base to rely on. By using a balanced combination of both implementations we

can “balance the impact of prior assumptions”. Seeger also highlights the fact that labelling the data is a taxing process, fortunately for us, existing data sets already exist with labelled and unlabelled images for flower species which will be useful for training. Therefore, a semi-supervised approach is feasible for the ML approach of this project.

4.1 Feature Extraction

Agarwal (2021) provides an introductory guide to feature extraction. They describe feature extraction as one of the two ways to reduce dimensionality with the other being feature selection. Extraction produces new features which are described as “linear combination of the existing features”. The process aims to use less features to encapsulate the same image information.

Tian (2013) conducts a review of image feature extraction techniques that are worth considering. They start with discussing extracting colour features such as histograms and colour “moments” from specific colour spaces such as RGB and HSV. The paper also compares different types of colour features, for example, histograms are simple to compute but are sensitive to noise. This feature will be important as flowers come in many different colours, but it cannot solely be relied on as different species can share similar colours. We can also extract information about the texture of an image, this is where we start thinking about analysing groups of pixels together. Texture in the context of images is a way to describe the perceived smoothness, roughness or bumpiness of an image through spatial variations in pixel intensity levels (MathWorks, n.d.). Lastly, the paper goes into depth about shape features and points to different sources that go into the subject with more depth, but to summarise, shape features are split into two broad categories of contour and region based. This is where the features are calculated from shape boundaries and image regions respectively. A simple example of shape feature is the circularity ratio where you measure how close a shape is to a circle by calculating the ratio of the area of a shape to the area of a circle with the same perimeter (Mingqiang, Kidiyo, Joseph, et al., 2008). Shape analysing will be very important in this project because flower shapes can differ greatly and can therefore serve as a way to easily differentiate between species.

4.2 Classification

Brownlee (2020) provides an easy-to-understand breakdown of classification within ML. They describe it as the process of assigning “a class label to example from the problem domain”. In our case, that means classifying a flower as species A as opposed to B, C or D. They also go into detail about the different classification methods such as:

- Binary classification e.g. it’s a flower or it’s not a flower.
- Multi-class classification, where we have more than two classes.

- Multi-label classification, this is where we have multiple predictions for classes based on a probability. This can be a path we could take if we wanted to produce multiple predictions for a flower species and then provide the likelihoods of each prediction to the user.

Next, we will look at classifiers, which help us carry out the classification stage. Fortunately, there is no shortage of types of classifiers in the ML space. Mohammed (2017) covers the most popular ones in good detail such as Naïve Bayes, k-Nearest Neighbour and Support Vector Machines (SVM). Starting with Naïve Bayes, this is a supervised classifier based on probability that assumes all attributes are independent:

$$P(c|E) = \frac{P(E|c)P(C)}{P(E)} \quad (1)$$

Where E is classified as the class $C = +$ if and only if

$$BC(E) = \frac{P(C = +|E)}{P(C = -|E)} \geq 1$$

BC is our Bayesian classifier, $+$ and $-$ are two separate classes (Zhang, 2004).

Zhang (2004) states that Naïve Bayes is superb in classification and demonstrates the classifier based version of it in Equation 1. They explore the optimal conditions of Naïve Bayes and propose that it is most optimal when the dependencies among attributes cancel out since Naïve Bayes works best when each attribute is independent, this is relevant as we want to explore optimisation in the project to ensure the highest level of efficiency when making the flower predictions.

Mohammed (2017) states that K-Nearest Neighbours (KNN) is one of the “simplest” of all the ML algorithms. Rosebrook (2016) discusses how to implement an image classifier using KNN where we can convert an image into a set of feature vectors on a graph, any new points get classified based on the k number of nearest neighbouring points, there’s no real learning in this process, just the calculation of where the nearest points are, based on (usually Euclidean) distance.

Noble (2006) describes SVM as a way to tackle binary classifications, which means in the context of flower classification it only really answers questions like “is it Flower A?”. They state that you would need to train multiple “one-versus-all” classifiers. For this project, that may not be appropriate if we want to support a large number of flower species.

ML techniques are certainly not useless and can still provide results, however, the research in the space has evolved to a new level, aiming to improve upon these traditional ML techniques in all evaluation categories. This is where Deep Learning (DL) comes in.

5 Deep Learning

DL will of course be our alternative approach to recognizing flower species. Where ML is basically our baseline, our DL implementation should hopefully highlight how much better it is compared to the ML approach.

5.1 Neurons and Perceptrons

Scarpino (2018) introduces the concept of Perceptrons in their book about using TensorFlow to implement DL. First, we must discuss neurons and how they relate to understanding the foundation of DL.




Figure 2: Simple diagram of a neuron (Scarpino, 2018).

What they choose to highlight in particular are three points that describe a neuron's functionality (and ultimately how it relates to perceptrons):

- A neuron receives one or more incoming signals and produces one outgoing signal.
- A neuron's output can serve as the input of another neuron.
- Every neuron has a threshold, and the neuron won't produce output until its electricity exceeds the threshold.

This page by Anon (n.d.) highlights a brief history of perceptrons, though it serves as a starting point to learn more about the concept. Perceptrons were coined by Frank Rosenblatt in 1962 (Rosenblatt, 1961), though his research is a bit outdated for our analysis, therefore here is a more modern version: Each input on the left is weighted and the summed within the circle node. If the summation meets a certain threshold, the output will be 1, if it doesn't meet the threshold then a 0 is outputted (Scarpino, 2018). Scarpino highlights some improvements to the model that was made including the weights that we discussed earlier as well as additional biases assigned with the incoming signals and an "activation function" that generates the output signal. Scarpino




Figure 3: Diagram of a perceptron (Anon, n.d.(b)).

goes further by linking activation functions directly with in built TensorFlow functions that carry out the same task. Making it quite useful to understand the link between the TensorFlow API and the underlying DL context. Once we start linking perceptrons and arranging them into layers we get a neural network as shown here:




Figure 4: Diagram of a layered network of perceptrons (Scarpino, 2018).

5.2 Convolutional Neural Networks

Bengio, Goodfellow and Courville (2015) and Scarpino (2018) go into detail about CNNs, Scarpino in particular is a useful source on how it works with image classification in TensorFlow. However, it's useful to have some sort of starting point for the subject. Saha (2018) highlights the key features of a CNN and their purposes such as the individual layers: convolutional (kernel), pooling and classification (see Figure 5).

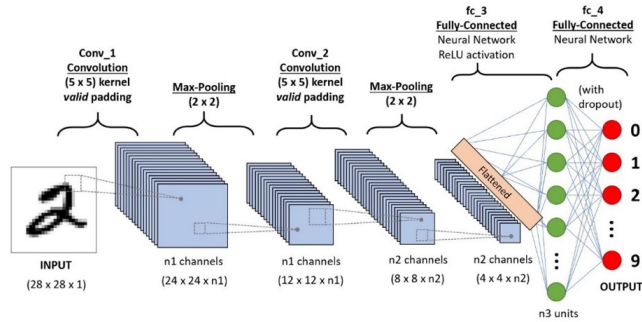


Figure 5: Example of a CNN process (Saha, 2018)

They also highlight a feature of CNNs that make images easier to process, where it reduces the size of images “into a form that is easier to process, without losing features which are critical for getting a good prediction”. This helps with our scalability approach when it comes to dataset sizes in particular. The ELI5 (Explain Like I’m 5) format is quite useful and allows us to highlight the key points of each layer (Saha, 2018):

- Convolution: Applying the kernel to extract high and low level features.
- Pooling: Reduces the spatial size of the output from the convolution. Decreases the “computational power” needed for data processing. Additionally, highlights features that are dominant.
- Classification: The pooling output is converted into column vectors and fed into a feed-forward neural network where the model is able to distinguish between features and classify them.

Saha (2018) also highlights that there are actually different implementations of CCNs, therefore they may not function in exactly the same format. Dive into Deep Learning (Anon, n.d.(a)) has a run down of multiple CNN types starting with LeNet-5 and more modern approaches like AlexNet, VGG, NiN, GoogLeNet, etc. We can also see how to implement them using TensorFlow which will prove useful when implementing our own solution for flowers.

Goodfellow, Bengio, and Courville (2015) go into detail about Deep Learning from the concept of perceptrons to modern implementations. Something they talk about that is interesting is the increasing data set and model sizes over time, which is quite applicable to the project since we are working with modern mobile hardware. They discuss how the increasing capabilities of computer hardware have led to the development of larger models and that neural networks tend to double in size roughly every 2.4 years. They predict that the trend will continue further on in the future. What is also relevant is the dataset sizes. Storing datasets take up storage space, they state that a deep learning algorithm (as of 2015) is stated to perform at acceptable levels with “around 5,000 labelled

examples per category, and will match or exceed human performance when trained with a dataset containing at least 10,000,000 labelled examples. That is of course, extremely large and is most definitely going to take up a lot of storage. Therefore, the book does re-iterate the earlier point of making use of unlabelled data like with semi-supervised learning. Goodfellow, Bengio, and Courville (2015) dig deep into the subject of DL and explain subjects from the applied maths to the modern practices of DL and the research in the field. This can prove useful in fully understanding the various processes in place including optimisations and ways to increase accuracy that we will need to consider when designing a DL model for the mobile application.

5.3 ML vs DL

DL is an obvious evolution from ML, but it is worth highlighting the key differences for clarity because ultimately this project will compare how ML and DL compete with each other. Kavlakoglu (2020) breaks down how DL is different from ML. They highlight that DL takes the initiative by automating feature extraction to lessen human intervention and that ML is more reliant on humans, where humans normally define the characteristics to look out for, as well as their priorities. DL is stated to “require more data points to improve its accuracy” compared to the ML counterpart.

Xin et al. (2018) goes into depth about the key differences when discussing approaches to ML and DL in the context of cybersecurity. However, the same reasoning can be applied in our situation. The key points they highlight are:

- Data dependencies: DL performs better with larger datasets like mentioned earlier as well as ML outperforming DL with smaller data sets.
- Hardware dependencies: DL requires a lot of matrix calculations and therefore a Graphical Processing Unit (GPU) can be used to optimise these processes. Note that mobile hardware do contain GPU hardware but they are not on the same scale as dedicated GPUs you find in PC hardware. Therefore, it will be interesting to see how DL fares against ML when we keep this hardware dependency in mind.
- Feature processing: Once again iterating on the point mentioned before, DL can extract features directly from the data and requires less human intervention.
- Execution time: DL algorithms take a lot longer to train compared to ML, this is dependant on the amount of data.
- Interpretability: Because of the complexity of DL it is hard to determine how a DL algorithm generated a result, whereas ML is more clearer.

I have summarised the key points, but they go into much more detail which could be helpful in the evaluation stage of comparing the two approaches of ML and DL.

6 Flower Classification

We will discuss further how flower classification is carried out including the use of ML and DL techniques, what features are extracted, the datasets that are used and the key challenges.

6.1 Existing Methods

Starting with what is known as the “Hello world” of ML, Iris flower classification serves as a simple and easy to understand project for developers to implement. The idea is to classify between three classes: Versicolor, Setosa and Virginica. There are many tutorials that can be followed online, this particular one by DataFlair (n.d.) provides additional background information about ML as well as how it will apply to the Iris project which is useful for our understanding. The tutorial uses the features of sepal length/width and petal length/width to determine the class of a flower. By using those inputs, they use a SVM to predict the species of a flower with 96% accuracy. I mentioned earlier that it may not work well for a large number of species, but it certainly works well for a small number of classes.

Nilsback and Zisserman (2008) demonstrate the effectiveness of a multiple kernel SVM on the oxford flowers 17 dataset. They manipulate the flower data to get key features such as the colour HSV values, the flower texture, shape and histogram of gradients (HOG) which “captures the more global spatial distribution of the flower” like the where the petals are arranged. They achieved an accuracy of around 88.3%. This is impressive considering the key challenges they highlight within flower classification. They state that flowers can share a lot of similarities between classes which can make it difficult to differentiate between species. Flowers are also “non-rigid objects” and therefore can appear in many different variations. Overall, they do a good job of explaining their reasoning for their dataset, citing the large variation of representations for each flower, and how they extract features from it, as well as how to build the classifier.

For deep learning, there is an extensive study that looks at using transfer learning, which is a technique of retraining TensorFlow models for different data sets. Xia, Xu, and Nan (2017) use the Inception-v3 model to train a classifier for the Oxford-17 and Oxford-102 flower datasets. They go into detail about the steps that took place to carry out the transfer learning as well as how to reconfigure the last layer of the network to only have 17 and 102 outputs for each dataset as it defaults to 1000. They found that the model for the for Oxford-17 and Oxford-102 datasets produced 95% and 94% accuracy respectively. This is very impressive performance, and their breakdown will be helpful when it comes to my own implementation. The paper really outlines the simplicity and flexibility of the Google’s TensorFlow, however, we will still need to investigate if these great results will translate to a mobile implementation as well.

6.2 Existing Apps

A large part of the project is to develop a fully functioning app that is developed to be more like a commercial product in addition to using deep learning techniques for the flower classification. This means developing features that will aid the user with using the app outside of main use case of recognising flowers. We will look at existing solutions that are already real products used by real people.

Pl@ntNet is a popular tool that has more than 10 million downloads on the Google Play Store alone (Google, 2022). It relies on volunteers to validate images and a search engine to identify them. Joly et al. (2015) go into detail about the overall experience of the app as well as provide insight into how it works.

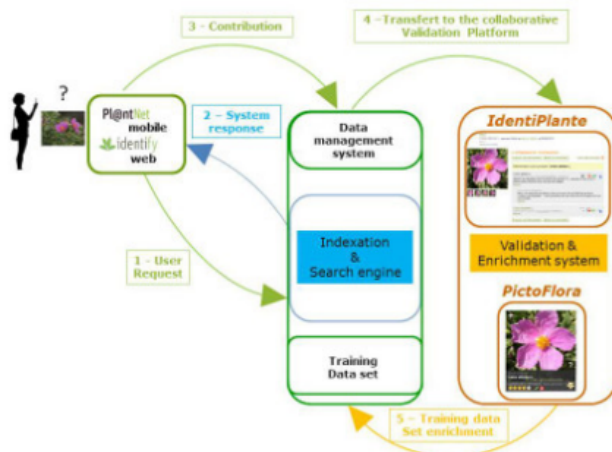


Figure 6: Diagram of Pl@ntnet user scenario (Joly et al., 2016).

Figure 6 shows clearly the type of system in place for the application. The user uses their device to query the search engine and get feedback, their image is also transferred the collaboration platform if they chose to. It can then be independently verified and added back to the training set. The search engine is then retrained on a nightly basis. The paper doesn't go into much more detail about the search engine itself apart from mentioning that progress in machine learning and computer vision should improve the performance of identification. Unfortunately, there aren't any new papers that provide a better look at the app, so it is hard to understand what changes have been made over the last several years as well as how what methods they use to build the search engine. There is however, a dataset now available for use that covers over a thousand plant species with over 300 thousand images (Garcin et al., 2021). This is of course out of the scope of the project as it doesn't strictly contain flowers, but it does go into detail about how to use the dataset as well as how to load the data and build a model with it.




Figure 7: Screenshot of PictureThis app (PictureThisAI, n.d.).

PictureThis is another plant identification app (shown above) that doesn't go into detail about how it works but does note that it requires an internet connection to function properly. This suggests that it must communicate with a server in order to generate a prediction for images. The app is very streamlined and has an easy-to-use UI that also contains useful features such as how to care for the plant and important information about it. My approach will be different in the sense that any identification process will be carried out on the device, however, it is still important to consider alternative methods and how they perform, so that we can compare approaches.

7 Evaluation

One of the key points of the project is having to evaluate the ML and DL approach, what we haven't discussed yet however, is how do we go about doing this?

Williams, Zander, and Armitage (2006) identify a process named "k-fold cross validation" where the data set is "divided into k subsets". One of the subsets is used to test the classifier and the rest (k-1) subsets is used as the training set. They use three performance metrics to test their ML systems: accuracy, precision and recall. Accuracy being the percentage of correct decisions over the total number of test instances. Confusion matrices can help us with representing accuracy by providing a "summary of prediction results" where we use the count of accurate and inaccurate predictions per class to show which particular classes

the classifier may be struggling with (Brownlee, 2020b). Precision and recall are a bit more complex. Fortunately, Shung (2018) demonstrates how these two differ to accuracy. Precision is the number of instances that are correctly determined over the total number of instances that are guessed, this is made up of correctly guessed instances as well as instances that are incorrectly guessed. Recall is the number of correctly predicted instances over the true number of instances in the class. In addition to these evaluation methods, Williams, Zander, and Armitage (2006) outline measuring CPU and memory usage. Fortunately, TensorFlow (lite) contains benchmarking tools for us to measure: Initialization time, inference time of warmup state/steady state, memory usage during initialization time and overall memory usage (Google, 2021). I will also assess real world speed and accuracy by analysing the app’s performance during development. The TensorFlow (lite) guide also contains tutorials on how to choose the best model for the task by comparing model size and accuracy of different models as well as the time it takes to make the prediction. The lite version is designed specifically for mobile and internet of things (IoT) hardware, so the additional tools will prove useful for the project later when we are at the evaluation stage.

Chockwanich and Visoottiviseth (2019) use the same evaluation methods outlined when comparing different DL models implemented in TensorFlow. They also look at CPU usage percentages and processing time. They were able to make a clear conclusion of which model is better by evaluating all factors. A term called f1-score was also mentioned in their analysis. Shung (2018) also explains the relevancy of f1 score, it is essentially a way to determine a “balance between precision and recall”. Korstanje (2021) discusses the F1 score and it’s purpose in providing a better accuracy statistic that accounts for “imbalanced data”, this is when you don’t have a good balance of data for each class and therefore the classifier makes inaccurate predictions heavily skewed towards classes you have more data for. The F1 score is calculated by:

$$\frac{2 * Precision * Recall}{Precision + Recall}$$

8 Final Thoughts

The review has really highlighted the progression ML and DL from the early concepts and how everything eventually fit together to form what we know today. The area will keep getting more exciting as we learn to optimise our current algorithms, come up with new ones and make use of advancing hardware capabilities. ML and DL is getting more and more accessible as manufacturers allow the use of their specialised hardware to developers who can make use of APIs built specifically for these tasks. Overall, we have built on the solid foundations of previous research and it’s interesting to see how the field develops in the future. The project hopes to aid the field by investigating the ML and DL approaches in the mobile format as well as explore why we see certain results

from the evaluation.

References

- Agarwal, D., 2021. *Guide for feature extraction techniques* [Online]. Available from: <https://www.analyticsvidhya.com/blog/2021/04/guide-for-feature-extraction-techniques/> [Accessed January 27, 2022].
- Anon, n.d.(a). *Convolutional neural networks* [Online]. Available from: https://d2l.ai/chapter_convolutional-neural-networks/lenet.html [Accessed December 2, 2021].
- Anon, n.d.(b). *History of the perceptron* [Online]. Available from: <https://home.csulb.edu/~cwallis/artificialn/History.htm> [Accessed November 30, 2021].
- Brownlee, J., 2020a. *4 types of classification tasks in machine learning* [Online]. Available from: <https://machinelearningmastery.com/types-of-classification-in-machine-learning/> [Accessed January 27, 2022].
- Brownlee, J., 2020b. *What is a confusion matrix in machine learning* [Online]. Available from: <https://machinelearningmastery.com/confusion-matrix-machine-learning/#:~:text=A%20confusion%20matrix%20is%20a%20summary%20of%20prediction%20results%20on,key%20to%20the%20confusion%20matrix> [Accessed March 9, 2022].
- Camstra, F. and Vinciarelli, A., n.d. *Machine learning for audio, image and video analysis: theory and applications*. 2nd ed. 2015, Advanced Information and Knowledge Processing. London: Springer London.
- Chockwanich, N. and Visoottiviseth, V., 2019. Intrusion detection by deep learning with tensorflow. *2019 21st international conference on advanced communication technology (icact)* [Online], pp.654–659. Available from: <https://doi.org/10.23919/ICACT.2019.8701969>.
- DataFlair, n.d. *Iris flower classification project using machine learning* [Online]. Available from: <https://data-flair.training/blogs/iris-flower-classification/#:~:text=Iris%20flower%20classification%20is%20a,%20%20Petal%20width> [Accessed January 28, 2022].
- Dietterich, T., 1995. Overfitting and undercomputing in machine learning. *Acm comput. surv.* [Online], 27(3) (), pp.326–327. Available from: <https://doi.org/10.1145/212094.212114>.
- Garcin, C., Joly, A., Bonnet, P., Affouard, A., Lombardo, J.-C., Chouet, M., Servajean, M., Lorieul, T., and Salmon, J., 2021. *Pl@ntnet-300k image dataset (v.1.1)* [Online]. Zenodo. Available from: <https://doi.org/10.5281/zenodo.5645731>.

- Goodfellow, I., Bengio, Y., and Courville, A., 2016. *Deep learning*. MIT press.
- Google, 2021. *Performance measurement* [Online]. Available from: <https://www.tensorflow.org/lite/performance/measurement> [Accessed December 2, 2021].
- Google, 2022. *Plantnet plant identification* [Online]. Available from: https://play.google.com/store/apps/details?id=org.plantnet&hl=en_GB&gl=US [Accessed March 7, 2022].
- GSMarena, 2021. *Gsm arena* [Online]. Available from: <https://www.gsmarena.com/> [Accessed November 30, 2021].
- Gupta, M., 2021. *Google tensor is a milestone for machine learning* [Online]. Available from: <https://blog.google/products/pixel/introducing-google-tensor/> [Accessed November 30, 2021].
- Joly, A., Bonnet, P., Goëau, H., Barbe, J., Selmi, S., Champ, J., Dufour-Kowalski, S., Affouard, A., Carré, J., Molino, J.-F., Boujemaa, N., and Barthélémy, D., 2016. A look inside the Pl@ntNet experience. *Multimedia Systems* [Online], 22(6), pp.751–766. Available from: <https://doi.org/10.1007/s00530-015-0462-9>.
- Kavlakoglu, E., 2020. *Ai vs. machine learning vs. deep learning vs. neural networks: what's the difference?* [Online]. Available from: <https://www.ibm.com/cloud/blog/ai-vs-machine-learning-vs-deep-learning-vs-neural-networks> [Accessed December 2, 2021].
- Korstanje, J., 2021. *The f1 score* [Online]. Available from: <https://towardsdatascience.com/the-f1-score-bec2bbc38aa6> [Accessed December 2, 2021].
- Kulendran, M., Lim, M., Laws, G., Chow, A., Nehme, J., Darzi, A., and Purkayastha, S., 2014. Surgical smartphone applications across different platforms: their evolution, uses, and users. *Surgical innovation*, 21(4), pp.427–440.
- MathWorks, n.d. *Texture analysis* [Online]. Available from: <https://www.mathworks.com/help/images/texture-analysis-1.html> [Accessed March 9, 2022].
- Mingqiang, Y., Kidiyo, K., Joseph, R., et al., 2008. A survey of shape feature extraction techniques. *Pattern recognition*, 15(7), pp.43–90.
- Mohammed, M., 2017. *Machine learning : algorithms and applications*. eng. 1st edition. Boca Raton: CRC Press.
- Nilsback, M.-E. and Zisserman, A., 2008. Automated flower classification over a large number of classes. *2008 sixth indian conference on computer vision, graphics image processing* [Online], pp.722–729. Available from: <https://doi.org/10.1109/ICVGIP.2008.47>.

- Noble, W.S., 2006. What is a support vector machine? eng. *Nature biotechnology*, 24(12), pp.1565–1567.
- O’Dea, S., 2021. *Global smartphone market share from 4th quarter 2009 to 3rd quarter 2021* [Online]. Available from: <https://www.statista.com/statistics/271496/global-market-share-held-by-smartphone-vendors-since-4th-quarter-2009/> [Accessed November 30, 2021].
- PictureThisAI, n.d. *Picturethis ai* [Online]. Available from: <https://www.picturethisai.com/> [Accessed March 7, 2022].
- Rosebrook, A., 2016. *K-nn classifier for image classification* [Online]. Available from: <https://www.pyimagesearch.com/2016/08/08/k-nn-classifier-for-image-classification/> [Accessed December 2, 2021].
- Rosenblatt, F., 1961. *Principles of neurodynamics. perceptrons and the theory of brain mechanisms*. (Technical report). Cornell Aeronautical Lab Inc Buffalo NY.
- Saha, S., 2018. *A comprehensive guide to convolutional neural networks — the eli5 way* [Online]. Available from: <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53> [Accessed December 2, 2021].
- Scarpino, M., 2018. *Tensorflow for dummies*. eng, For dummies. Newark: Wiley.
- Seeger, M., 2000. Learning with labeled and unlabeled data.
- Shung, K.P., 2018. *Accuracy, precision, recall or f1?* [Online]. Available from: <https://towardsdatascience.com/accuracy-precision-recall-or-f1-331fb37c5cb9> [Accessed December 2, 2021].
- Szeliski, R., 2011. *Computer vision : algorithms and applications*. eng. 1st ed. 2011., Texts in Computer Science. London: Springer London : Imprint: Springer.
- Three, 2021. *Samsung galaxy s21 ultra 5g* [Online]. Available from: <http://www.three.co.uk/samsung/galaxy-s21-ultra-5g?colour=phantom%20black&memory=128&paym=true> [Accessed December 2, 2021].
- Tian, D., 2013. A review on image feature extraction and representation techniques. *International journal of multimedia and ubiquitous engineering*, 8 (), pp.385–395.
- Triggs, R. and Simons, H., 2021. *Google tensor vs snapdragon 888 series: how the pixel 6’s chip shapes up* [Online]. Available from: <https://www.androidauthority.com/google-tensor-vs-snapdragon-888-3025332/> [Accessed November 30, 2021].

- Turk, M.A. and Pentland, A.P., 1991. Face recognition using eigenfaces. IEEE Computer Society, pp.586–587.
- Williams, N., Zander, S., and Armitage, G., 2006. A preliminary performance comparison of five machine learning algorithms for practical ip traffic flow classification. *Sigcomm comput. commun. rev.* [Online], 36(5) (), pp.5–16. Available from: <https://doi.org/10.1145/1163593.1163596>.
- Xia, X., Xu, C., and Nan, B., 2017. Inception-v3 for flower classification. *2017 2nd international conference on image, vision and computing (icivc)* [Online], pp.783–787. Available from: <https://doi.org/10.1109/ICIVC.2017.7984661>.
- Xin, Y., Kong, L., Liu, Z., Chen, Y., Li, Y., Zhu, H., Gao, M., Hou, H., and Wang, C., 2018. Machine learning and deep learning methods for cybersecurity. *Ieee access* [Online], 6, pp.35365–35381. Available from: <https://doi.org/10.1109/ACCESS.2018.2836950>.
- Zhang, H., 2004. The optimality of naive bayes. *Aa*, 1(2), p.3.
- Zhu, X.J., 2005. Semi-supervised learning literature survey.